



C  
h  
e  
f

# Agenda



# Topics

- Overview of Chef
- Workstation Setup
- Node Setup
- Chef Resources and Recipes
- Introducing the Node object
- Setting Node attributes
- Roles
- Community Cookbooks

# Lesson Objectives

- After completing the lesson, you will be able to
  - Describe how Chef thinks about Infrastructure Automation
  - Define the following terms:
    - Resource
    - Recipe
    - Node
    - Run List
    - Search

# Complexity

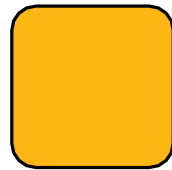


<http://www.flickr.com/photos/michaelheiss/3090102907/>

# Items of Manipulation (Resources)

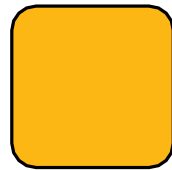
- Networking
- Files
- Directories
- Symlinks
- Mounts
- Registry Keys
- Powershell Scripts
- Users
- Groups
- Packages
- Services
- Filesystems

# A tale of growth...

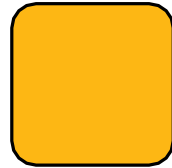


Application  
n

# Add a database



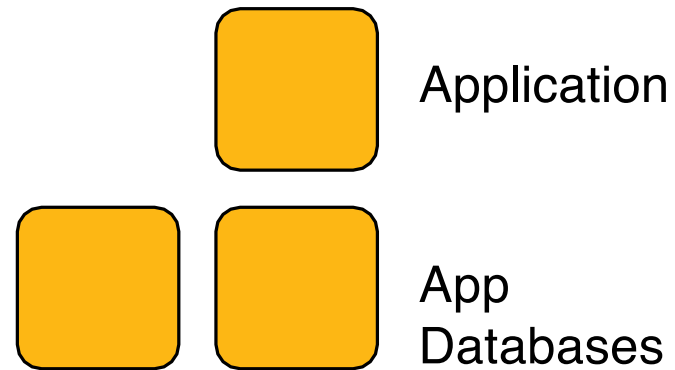
Application



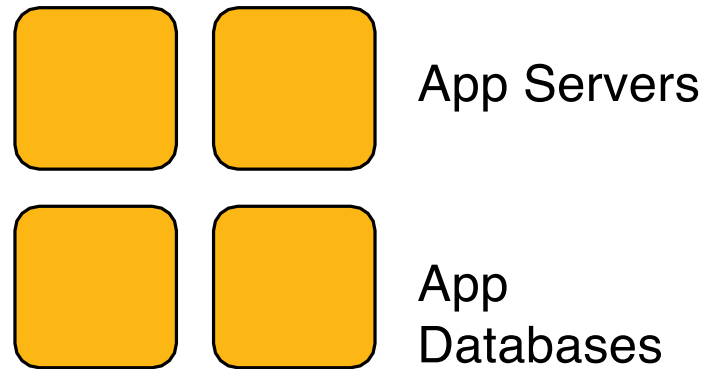
Application  
Database



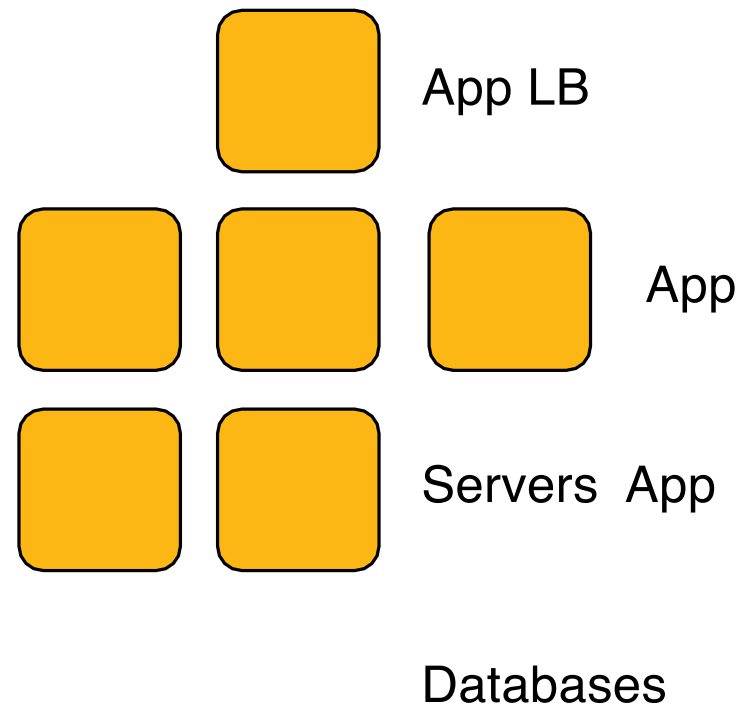
# Make database redundant



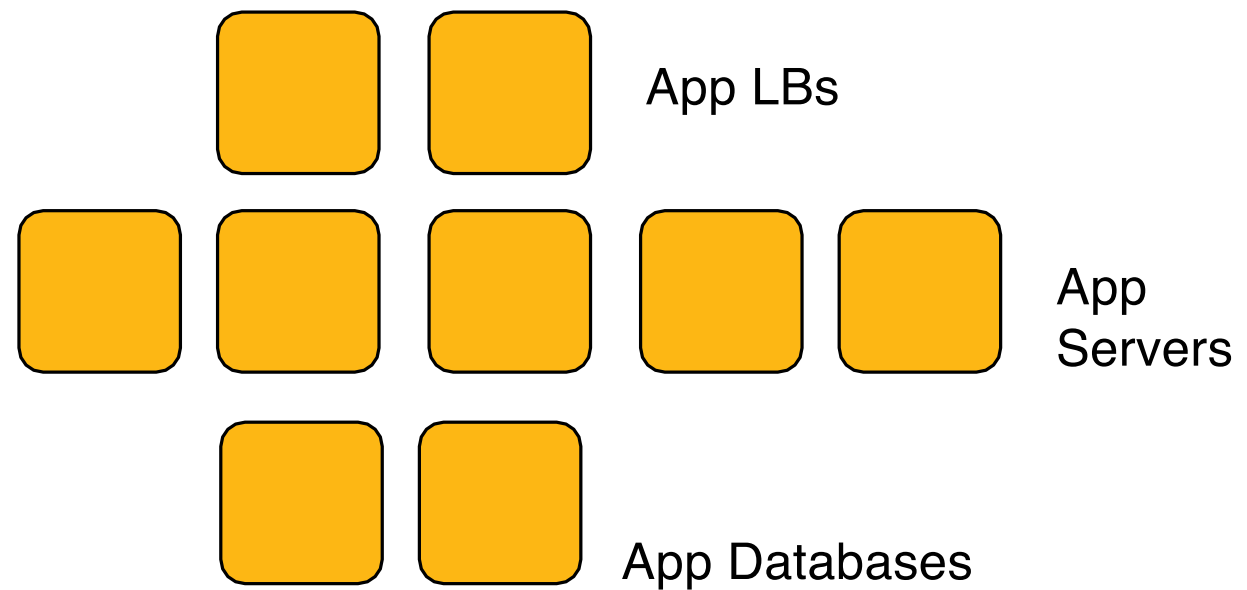
# Application server redundancy



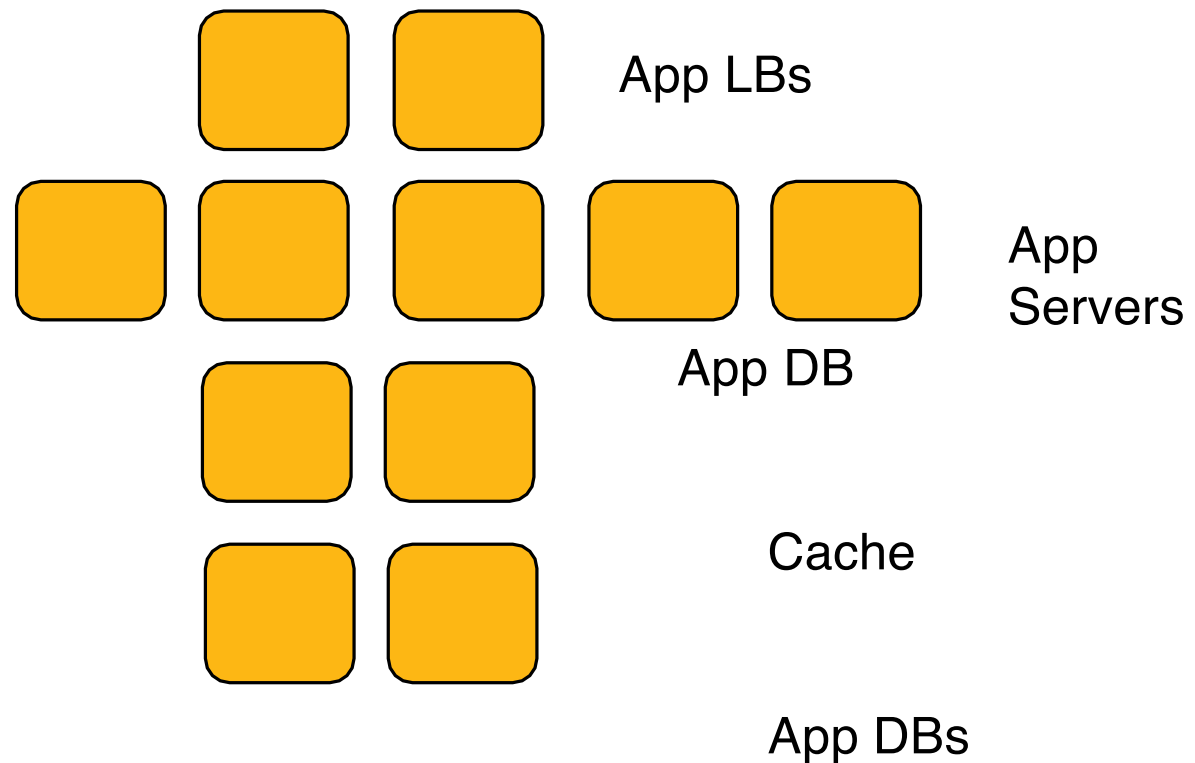
# Add a load balancer



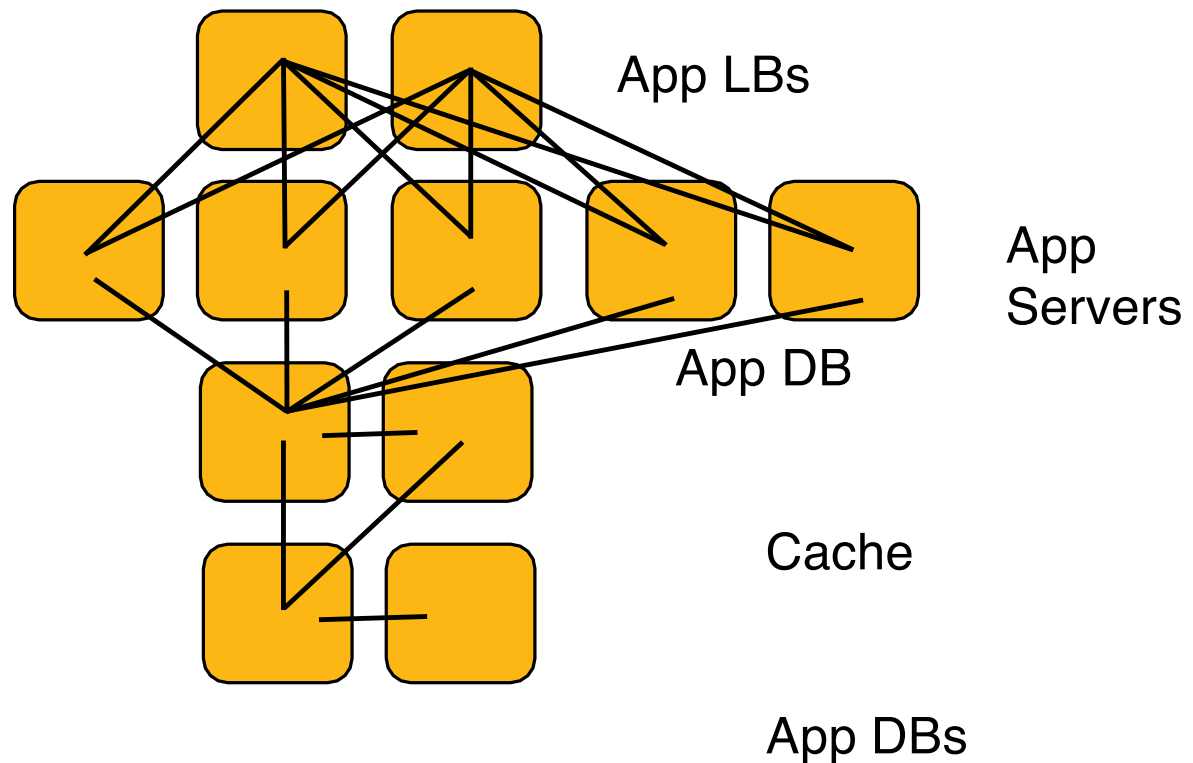
# WebScale!



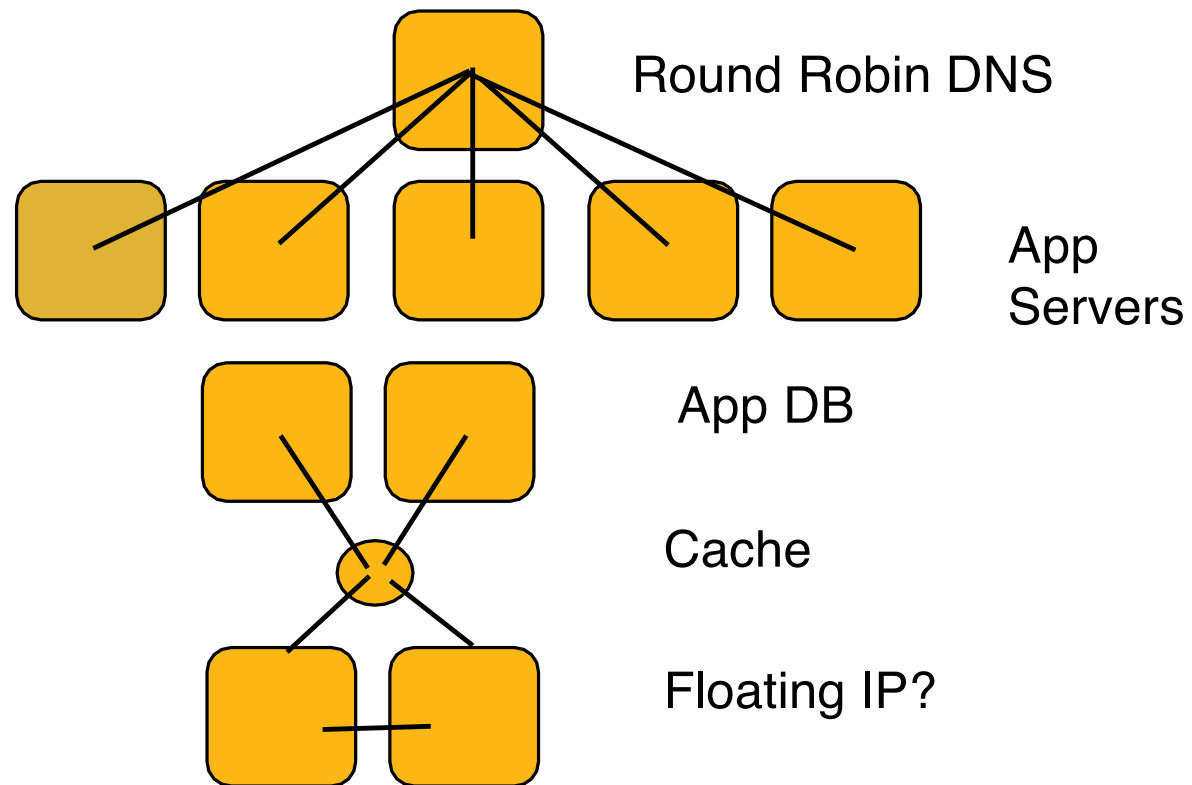
# Now we need a caching layer



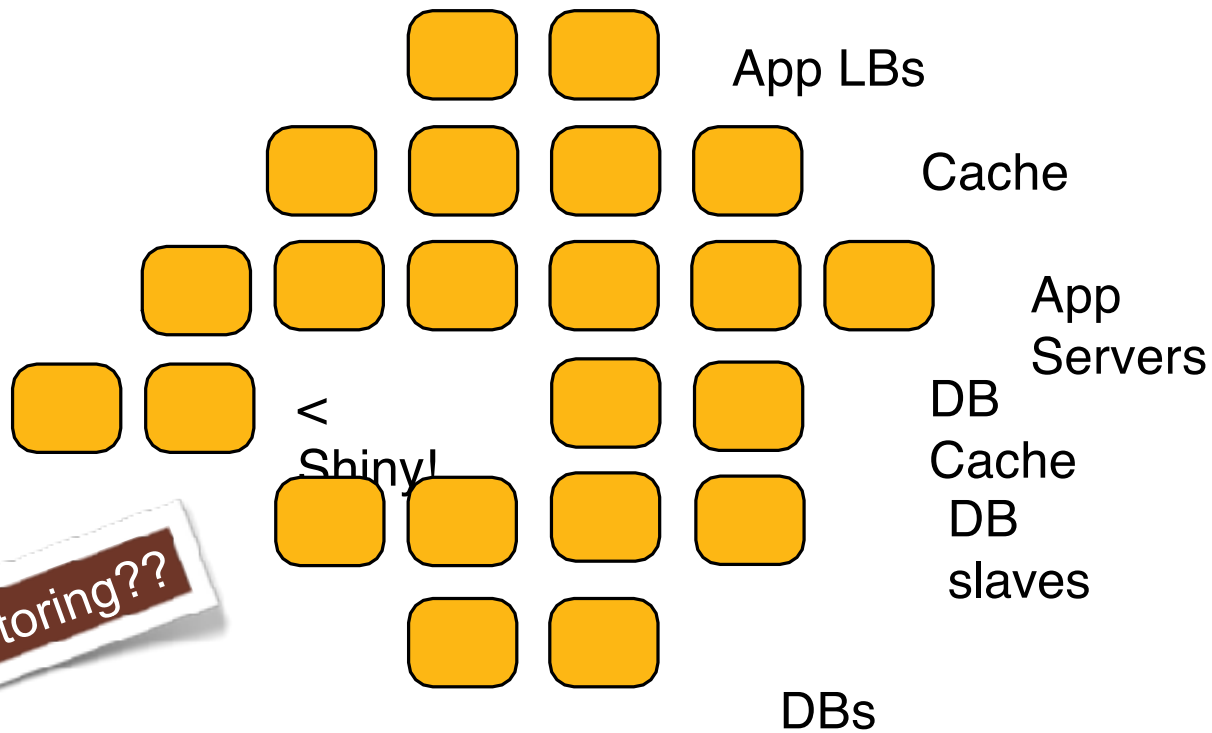
# Infrastructure has a Topology



# Your Infrastructure is a Snowflake

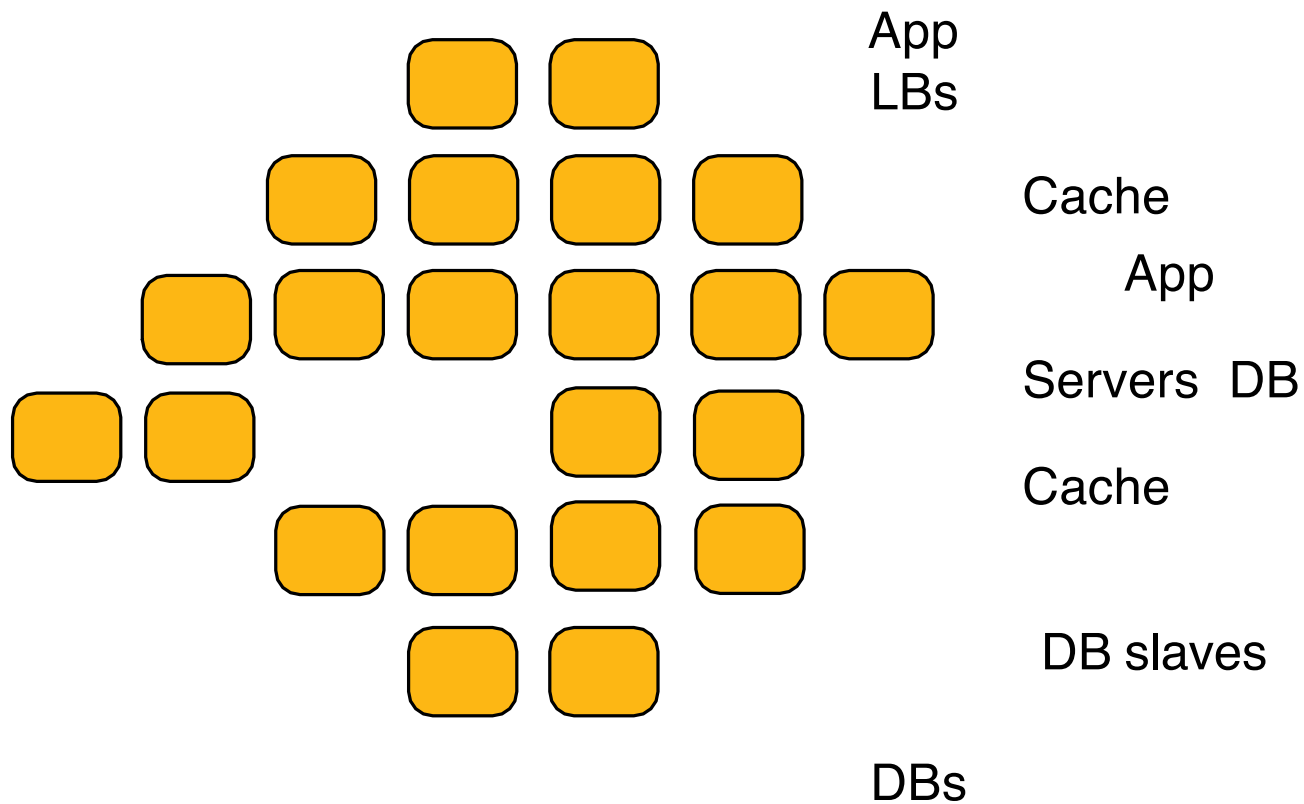


# Complexity Increases Quickly

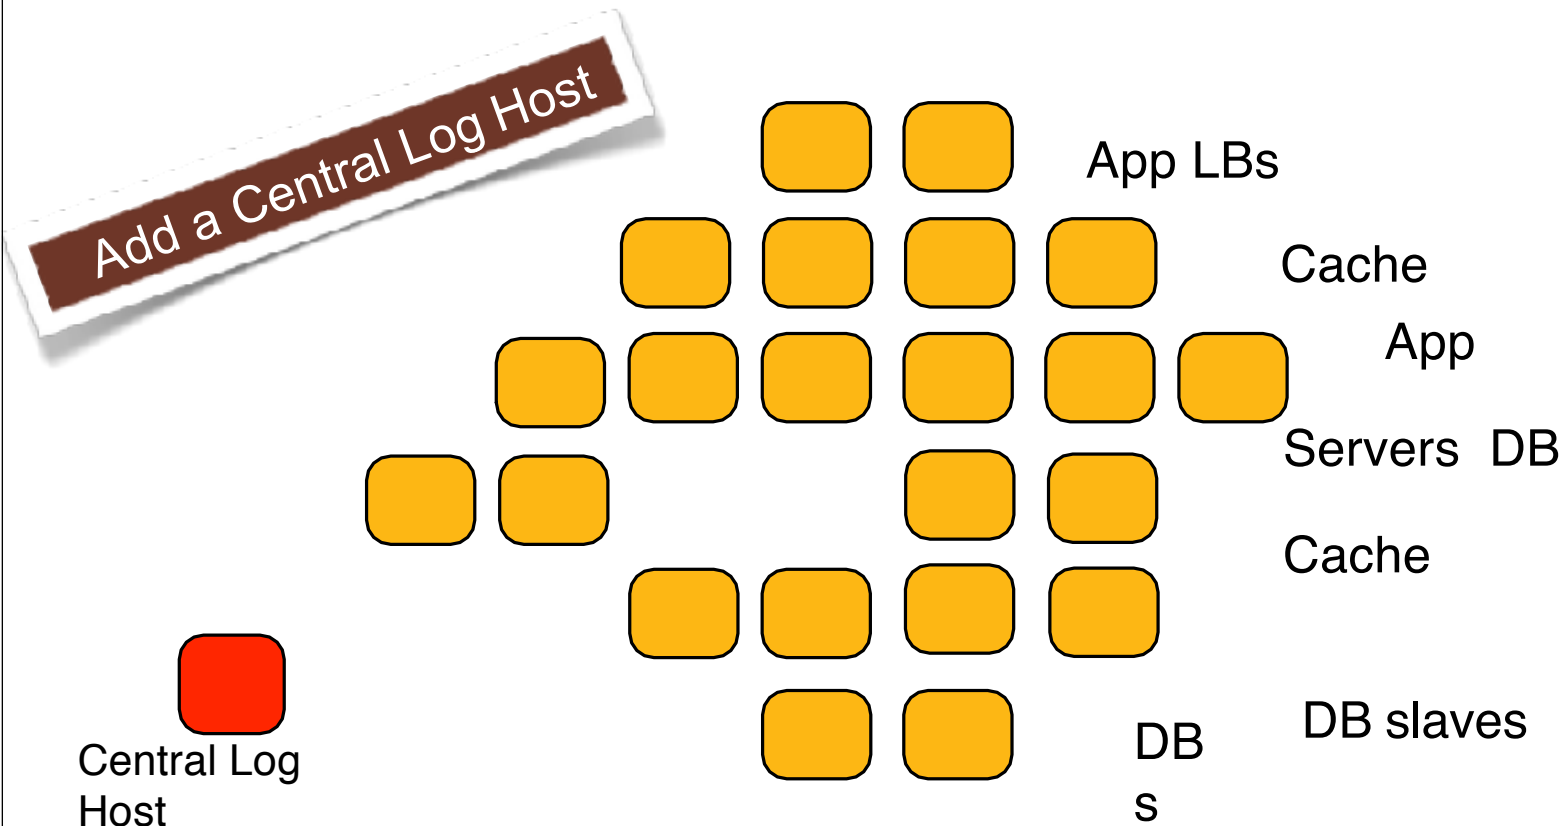




# ...and change happens!




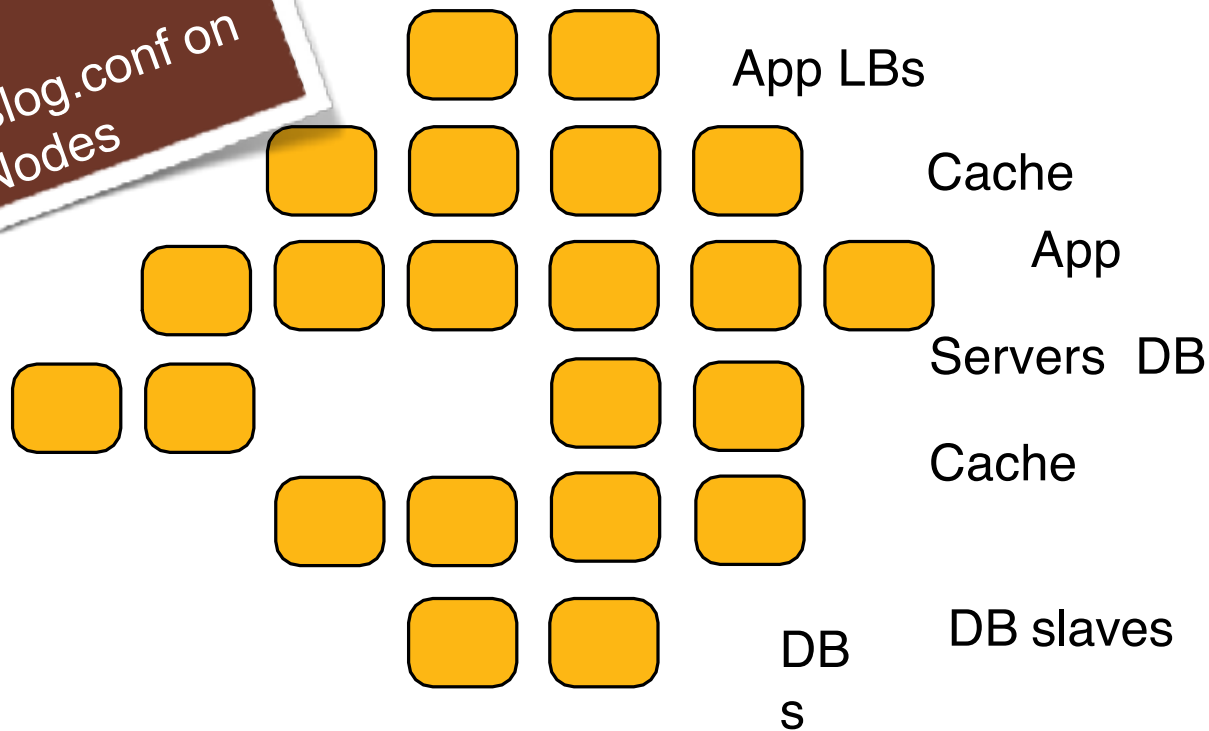
# ...and change happens!



# ...and change happens!

Add a Central Log Host  
Update syslog.conf on  
all Nodes

  
Central Log  
Host



# Chef Solves This Problem

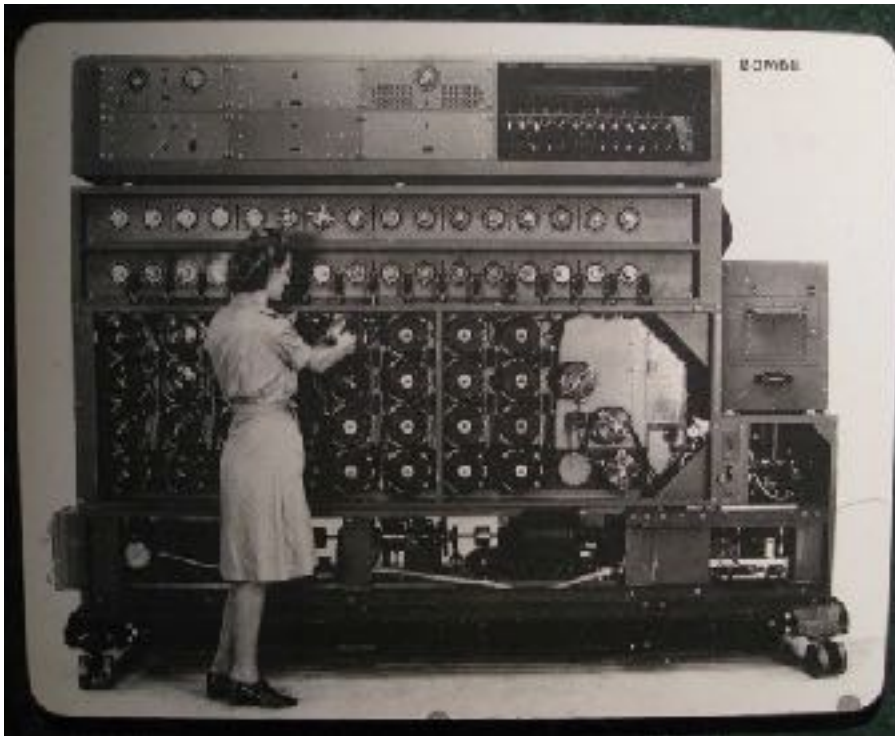


**CHEF**™  
GETCHEF.COM

- But you already guessed that, didn't you?

# Chef is Infrastructure as Code

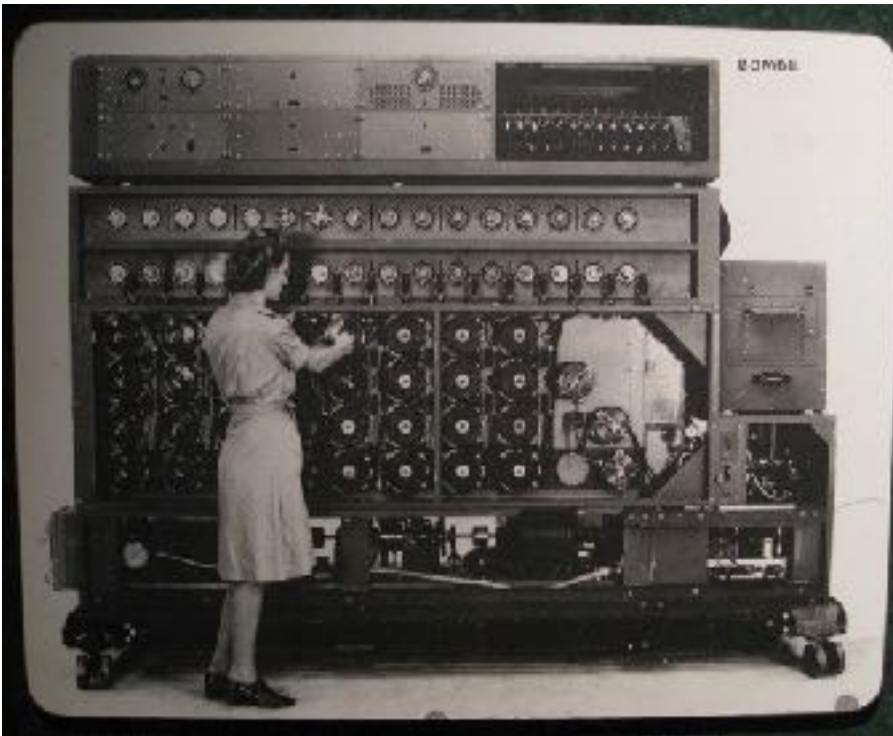
- Programmatically provision and configure components



<http://www.flickr.com/photos/louisb/4555295187/>

# Chef is Infrastructure as Code

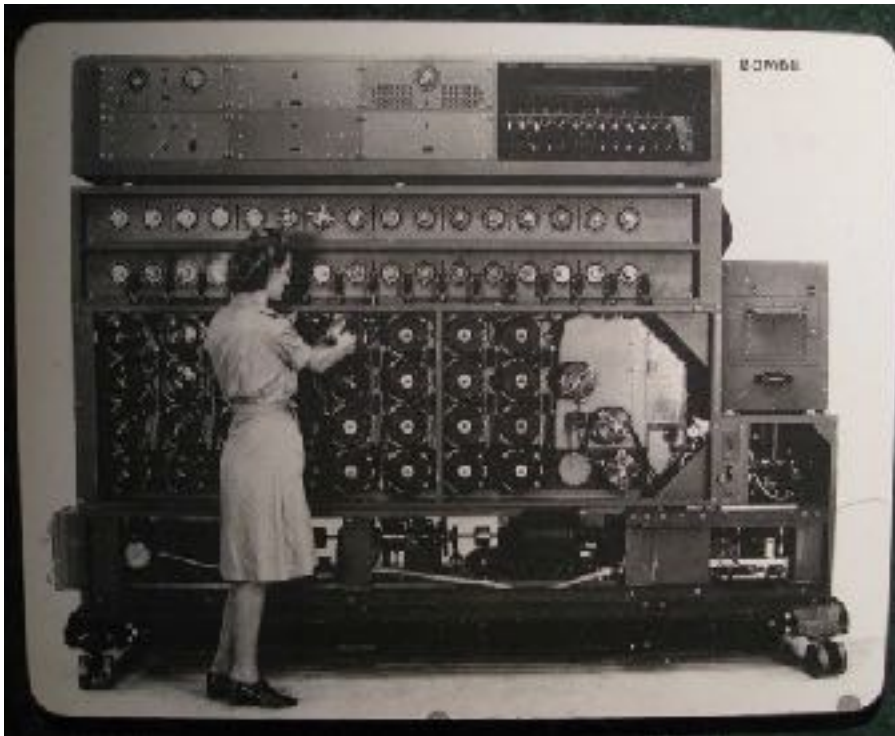
- Treat like any other code base



<http://www.flickr.com/photos/louisb/4555295187/>

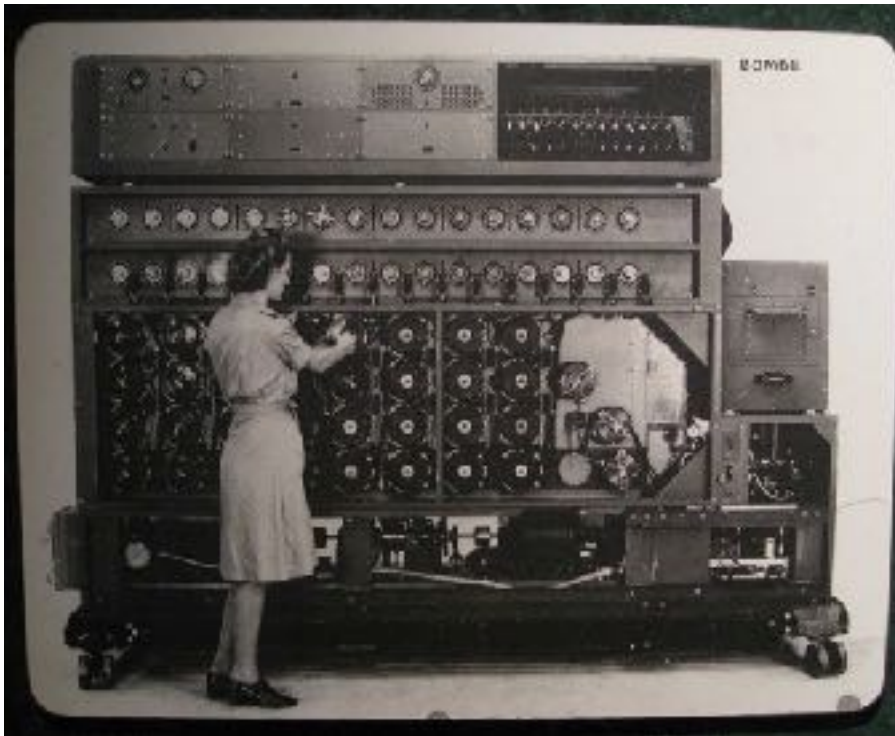
# Chef is Infrastructure as Code

- Reconstruct business from **code repository**, **data backup**, and **compute resources**



<http://www.flickr.com/photos/louisb/4555295187/>

# Chef is Infrastructure as Code



<http://www.flickr.com/photos/louisb/4555295187/>

- Programmatically provision and configure components
- Treat like any other code base
- Reconstruct business from **code repository**, **data backup**, and **compute resources**



# Configuration Code

- Chef ensures each Node complies with the policy
- Policy is determined by the configurations in each Node's run list
- Reduce management complexity through abstraction
- Store the configuration of your infrastructure in version control

# Declarative Interface to Resources

- You define the policy in your Chef configuration
- Your policy states what state each resource should be in, but not how to get there
- Chef-client will pull the policy from the Chef Server and enforce the policy on the Node

# Managing Complexity

- Resources
- Recipes
- Nodes
- Search

# Resources

- A Resource represents a piece of the system and its desired state
  - A package that should be installed
  - A service that should be running
  - A file that should be generated
  - A cron job that should be configured
  - A user that should be managed
  - and more

# Resources in Recipes

- Resources are the fundamental building blocks of Chef configuration
- Resources are gathered into Recipes
- Recipes ensure the system is in the desired state

# Recipes

- Configuration files that describe resources and their desired state
- Recipes can:
  - Install and configure software components
  - Manage files
  - Deploy applications
  - Execute other recipes
  - and more

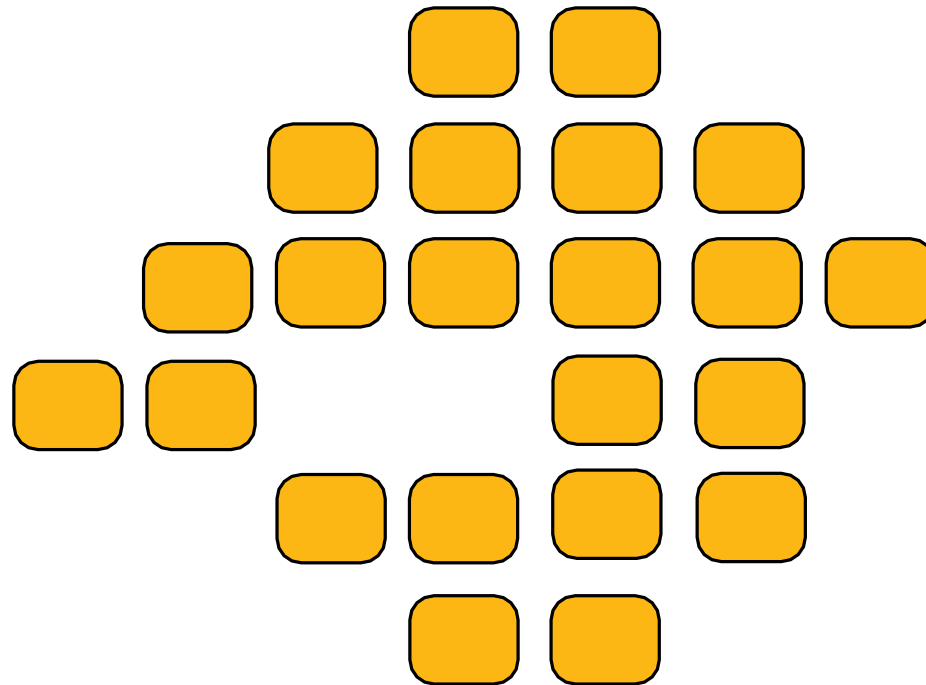
# Example Recipe

```
package "httpd" do
  action :start
end
```

```
template "/etc/httpd/conf/httpd.conf" do
  source "httpd.conf.erb"
  owner "root"
  group "root"
  mode "0644"
  variables(:allow_override => "All")
  notifies :reload, "service[httpd]"
end
```

```
service "httpd" do
  action [:enable, :start]
  supports :reload => true
end
```

# Nodes

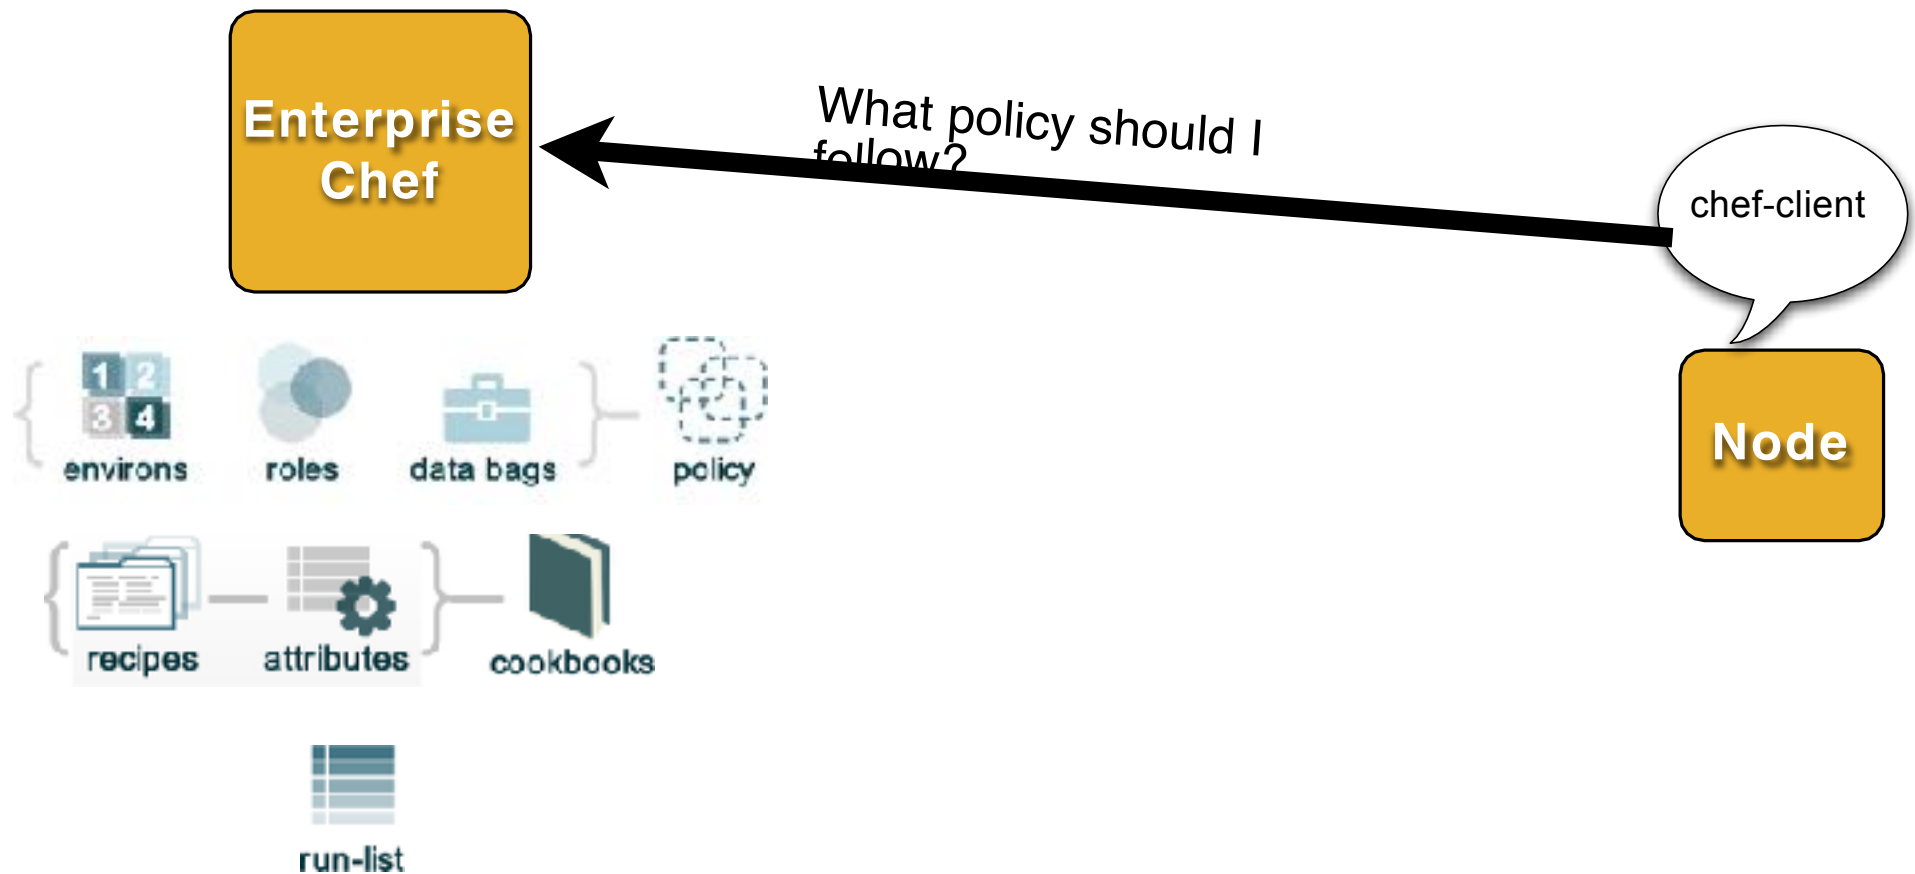




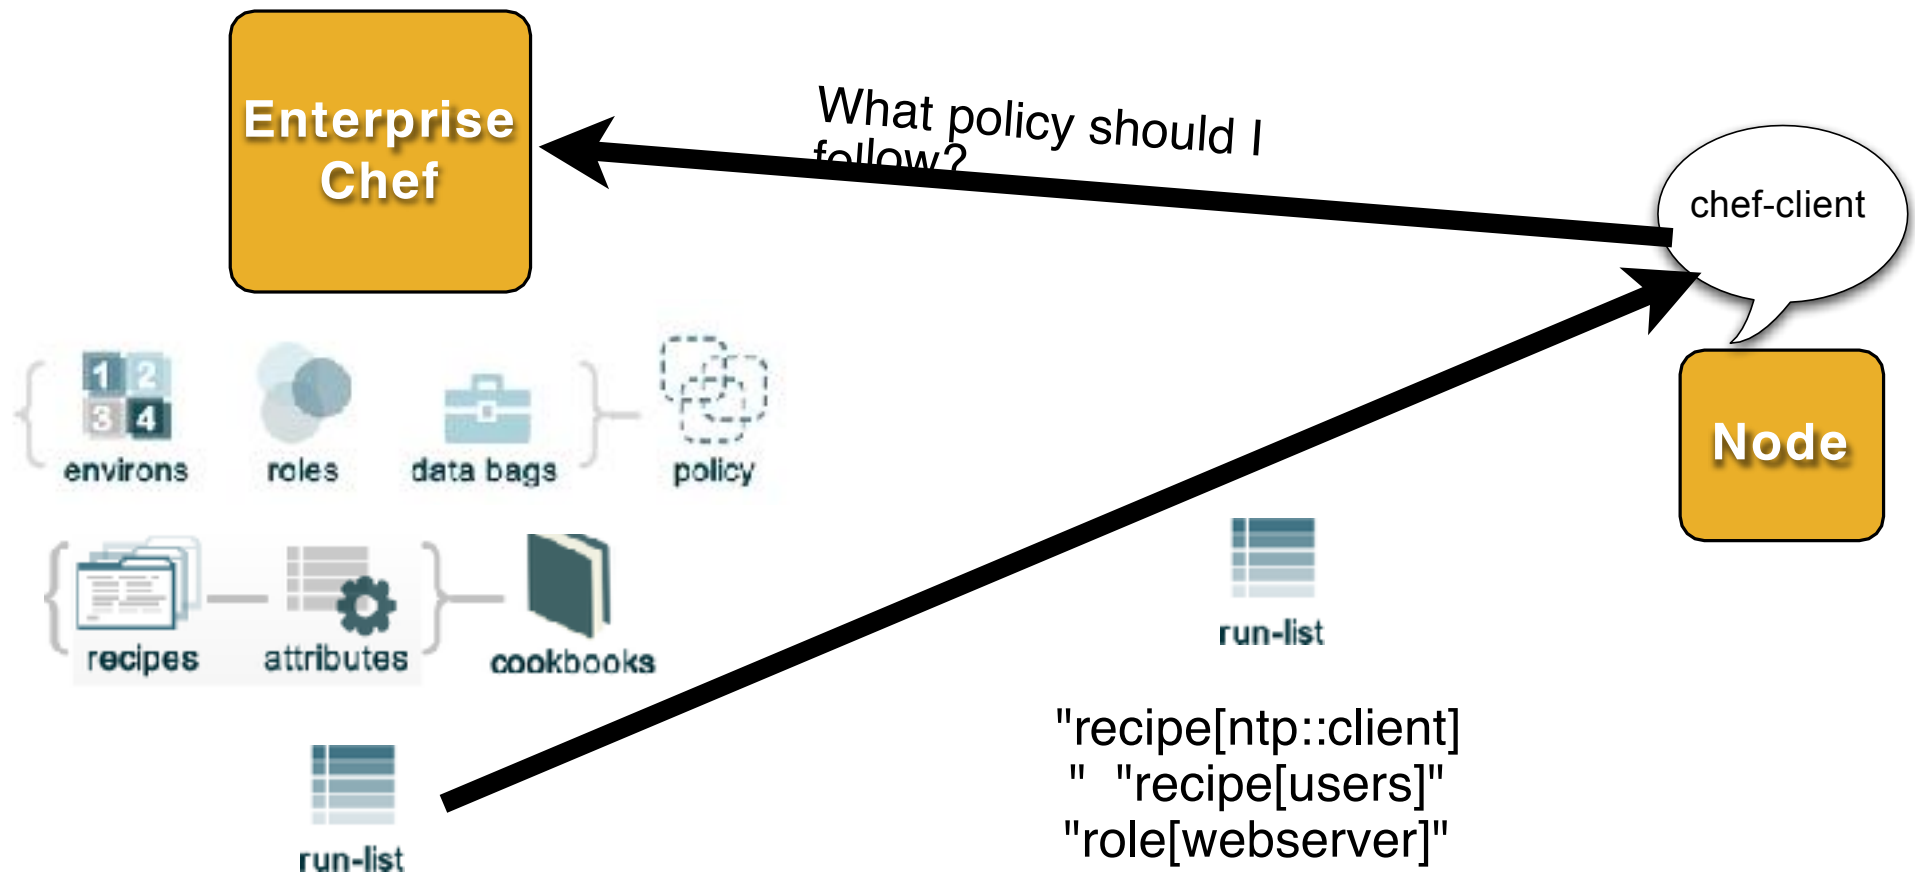
# Nodes Adhere to Policy

- The chef-client application runs on each node, which
  - Gathers the current system configuration of the node
  - Downloads the desired system configuration policies from the Chef server for that node
  - Configures the node such that it adheres to those policies

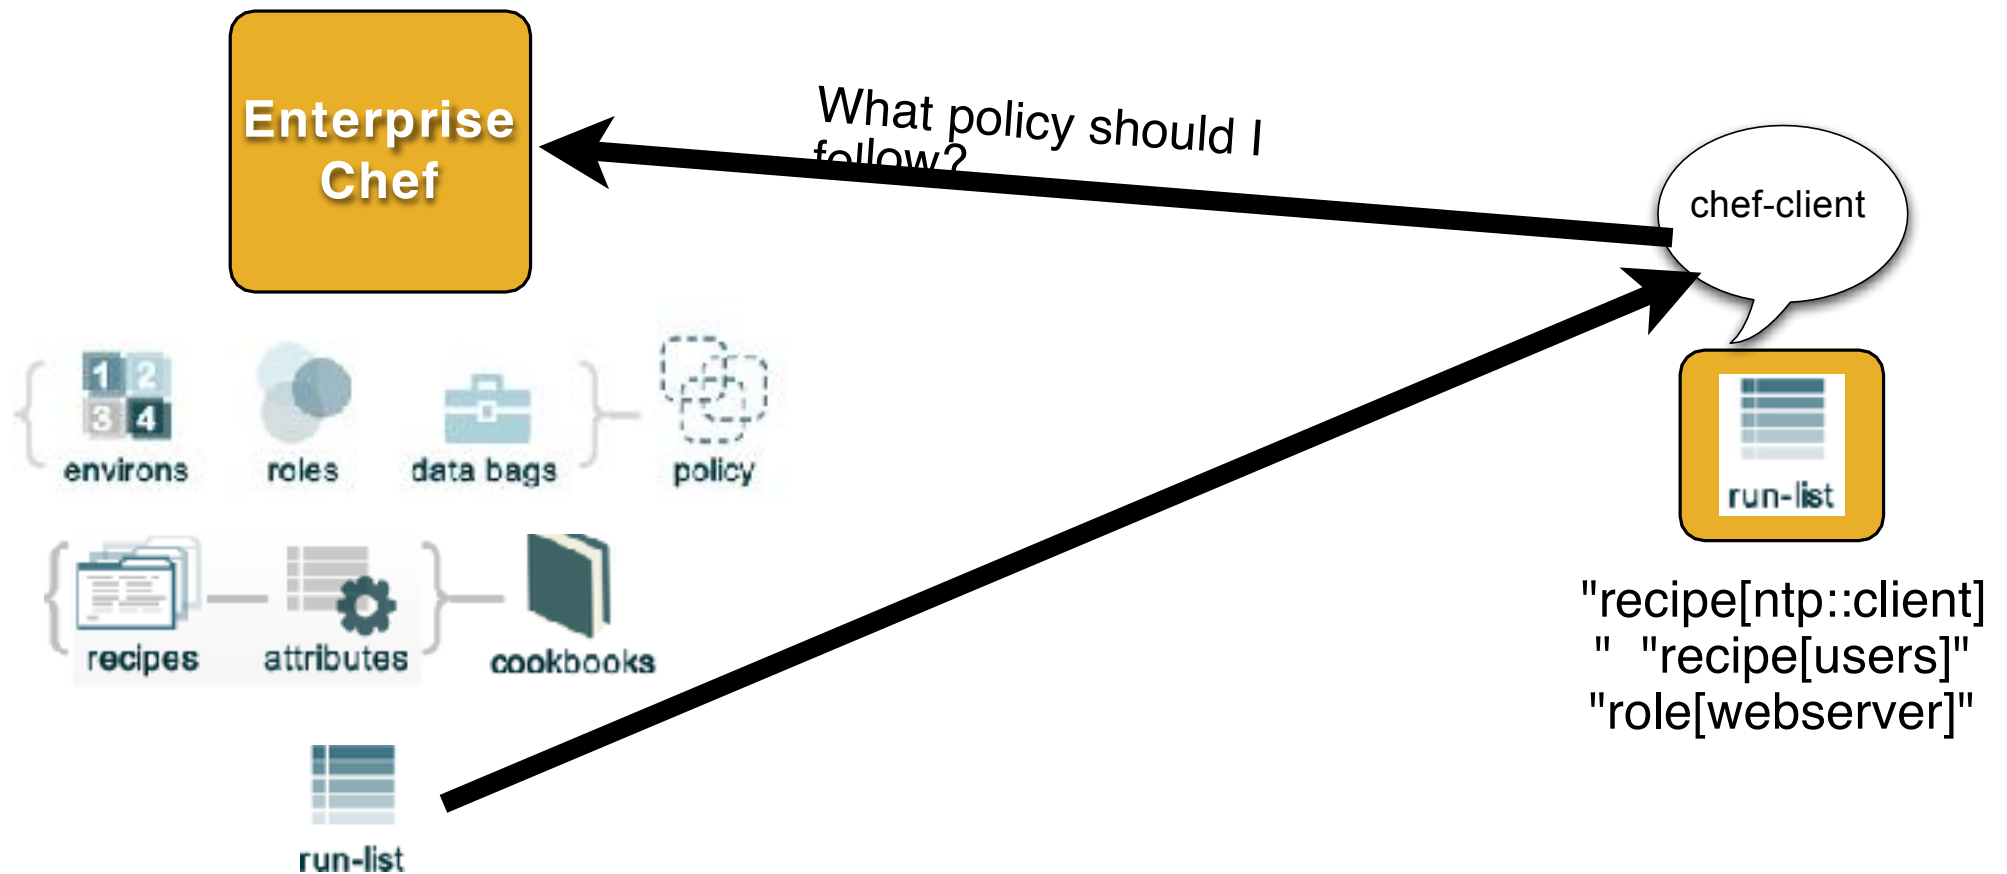
# Run List



# Run List



# Run List



# Run List Specifies Policy

- The Run List is an ordered collection of policies that the Node should follow
- Chef-client obtains the Run List from the Chef Server
- Chef-client ensures the Node complies with the policy in the Run List

# Search

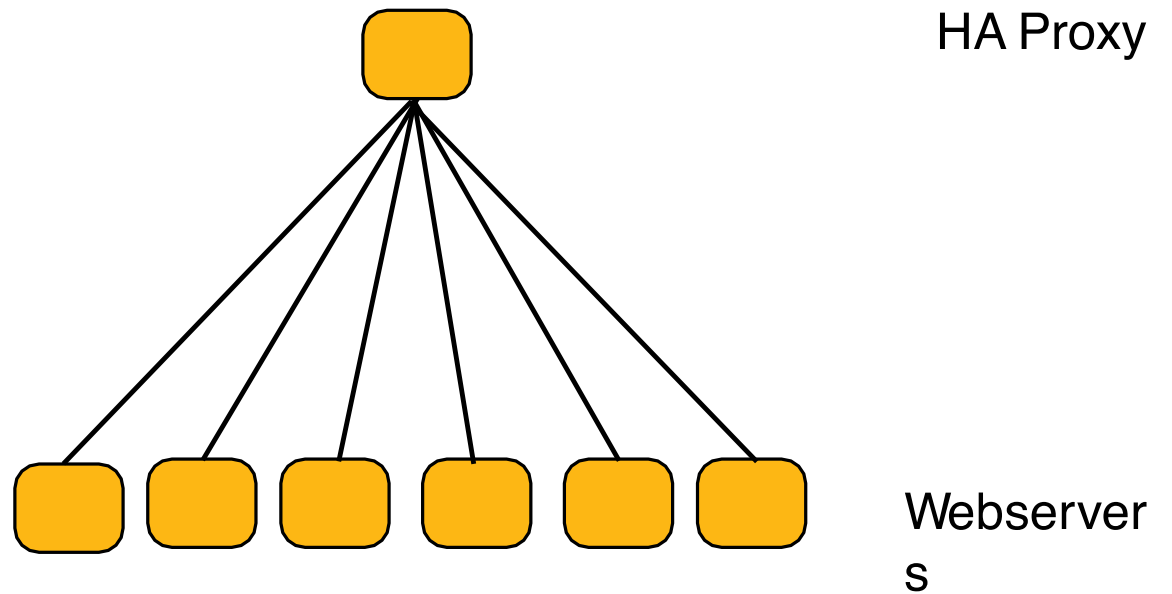
- Search for nodes with Roles
- Find Topology Data
- IP addresses
- Hostnames
- FQDNs

# Search for Nodes

```
pool_members = search("node", "role:webserver")

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy-app_lb.cfg.erb"
  owner  "root"
  group  "root"
  mode   0644
  variables :pool_members => pool_members.uniq
  notifies :restart, "service[haproxy]"
end
```

# HAProxy Configuration





# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```

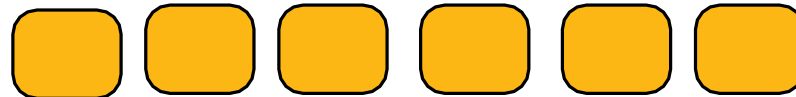
Enterprise  
Chef



HA Proxy



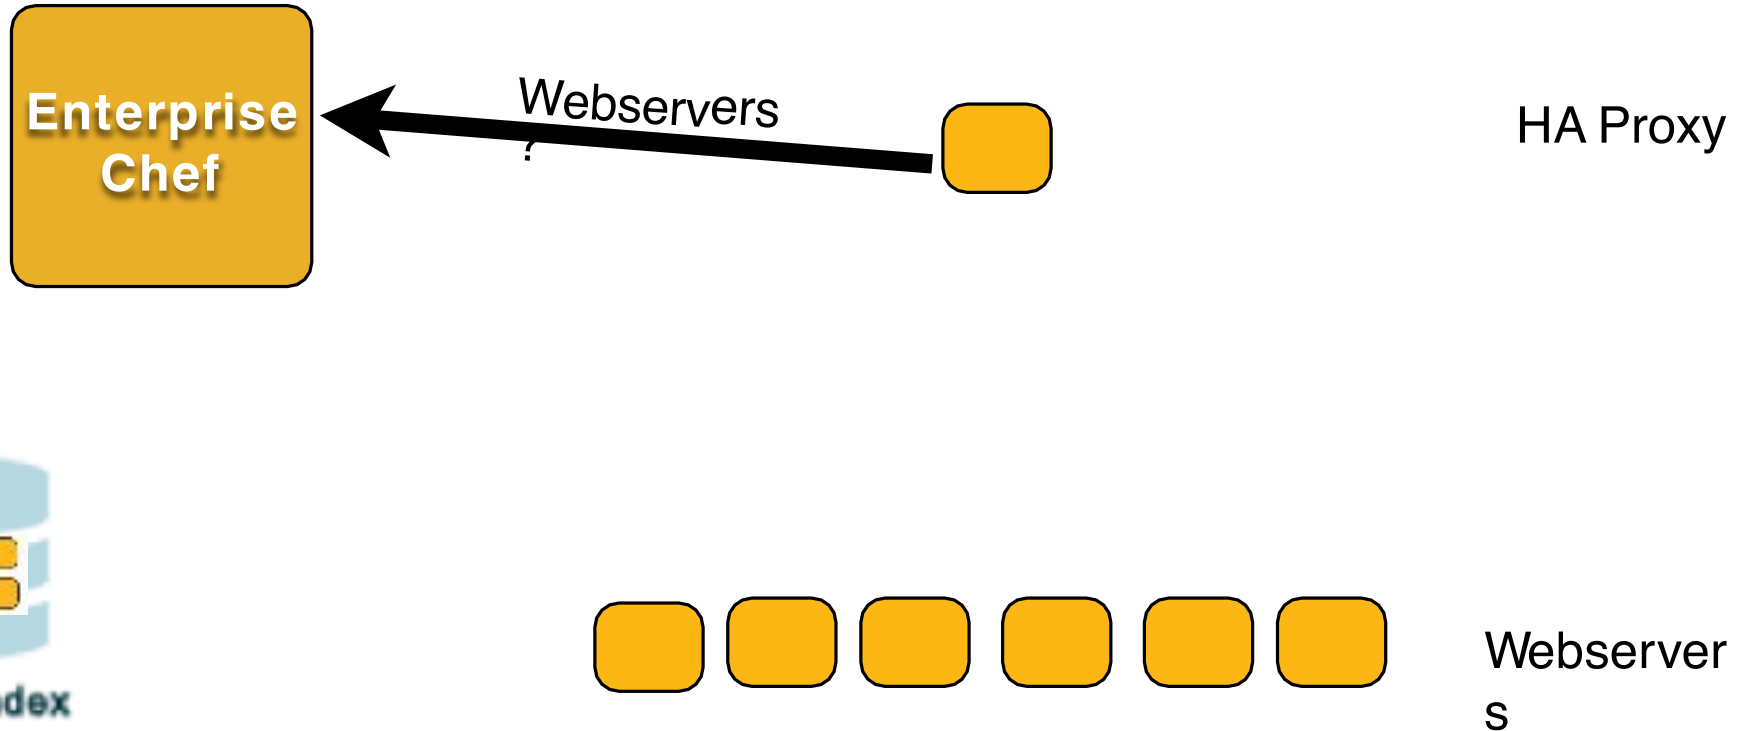
Search Index



Webserver  
s

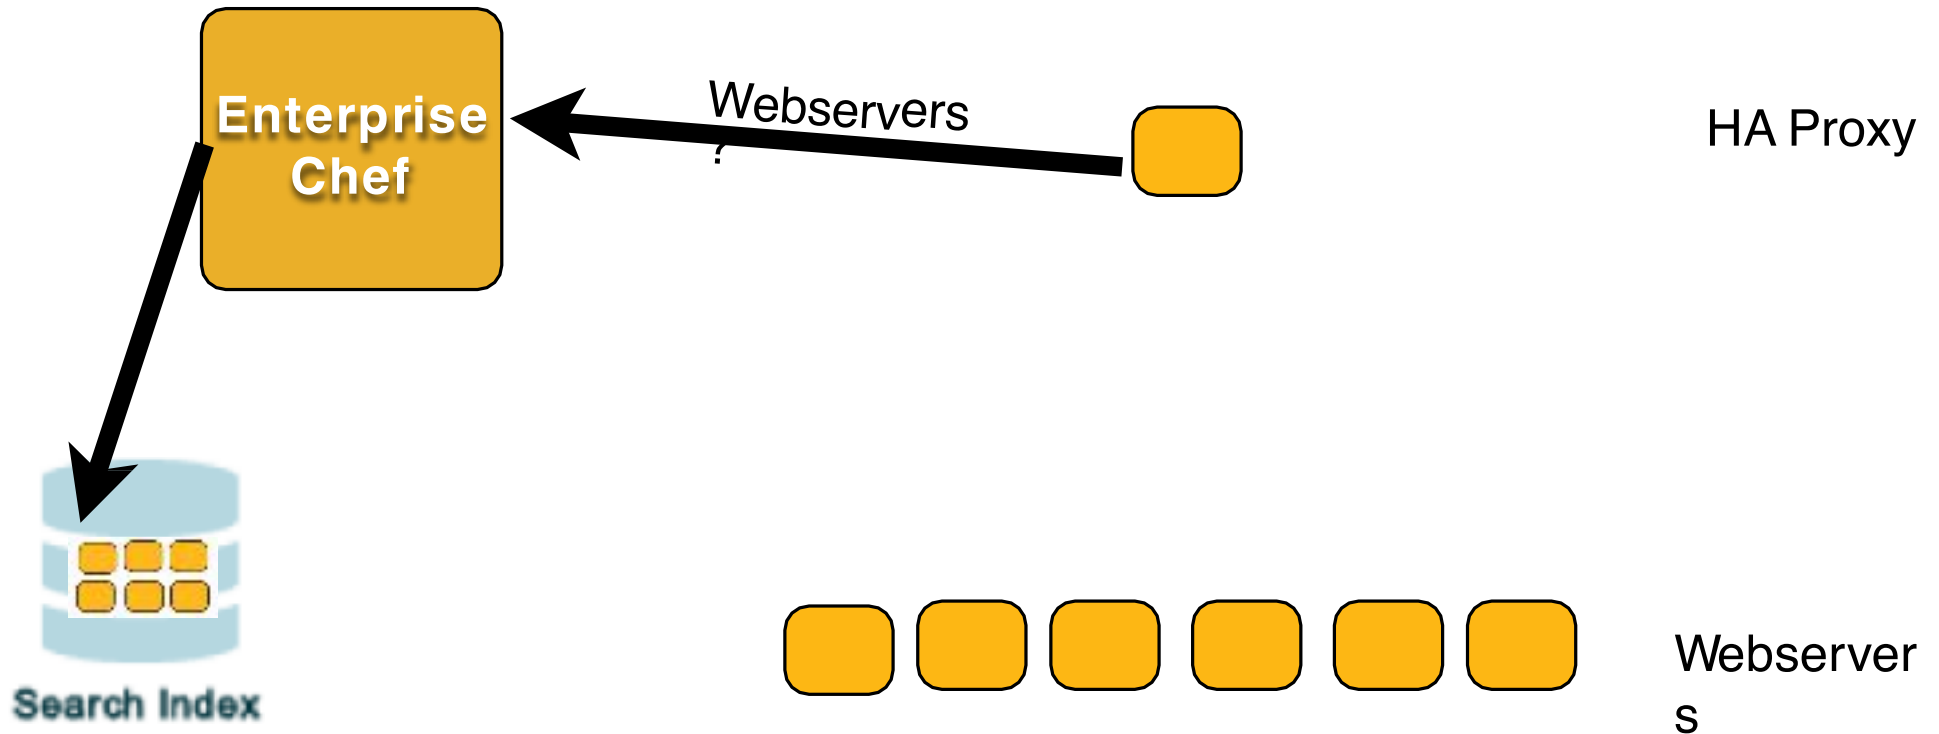
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



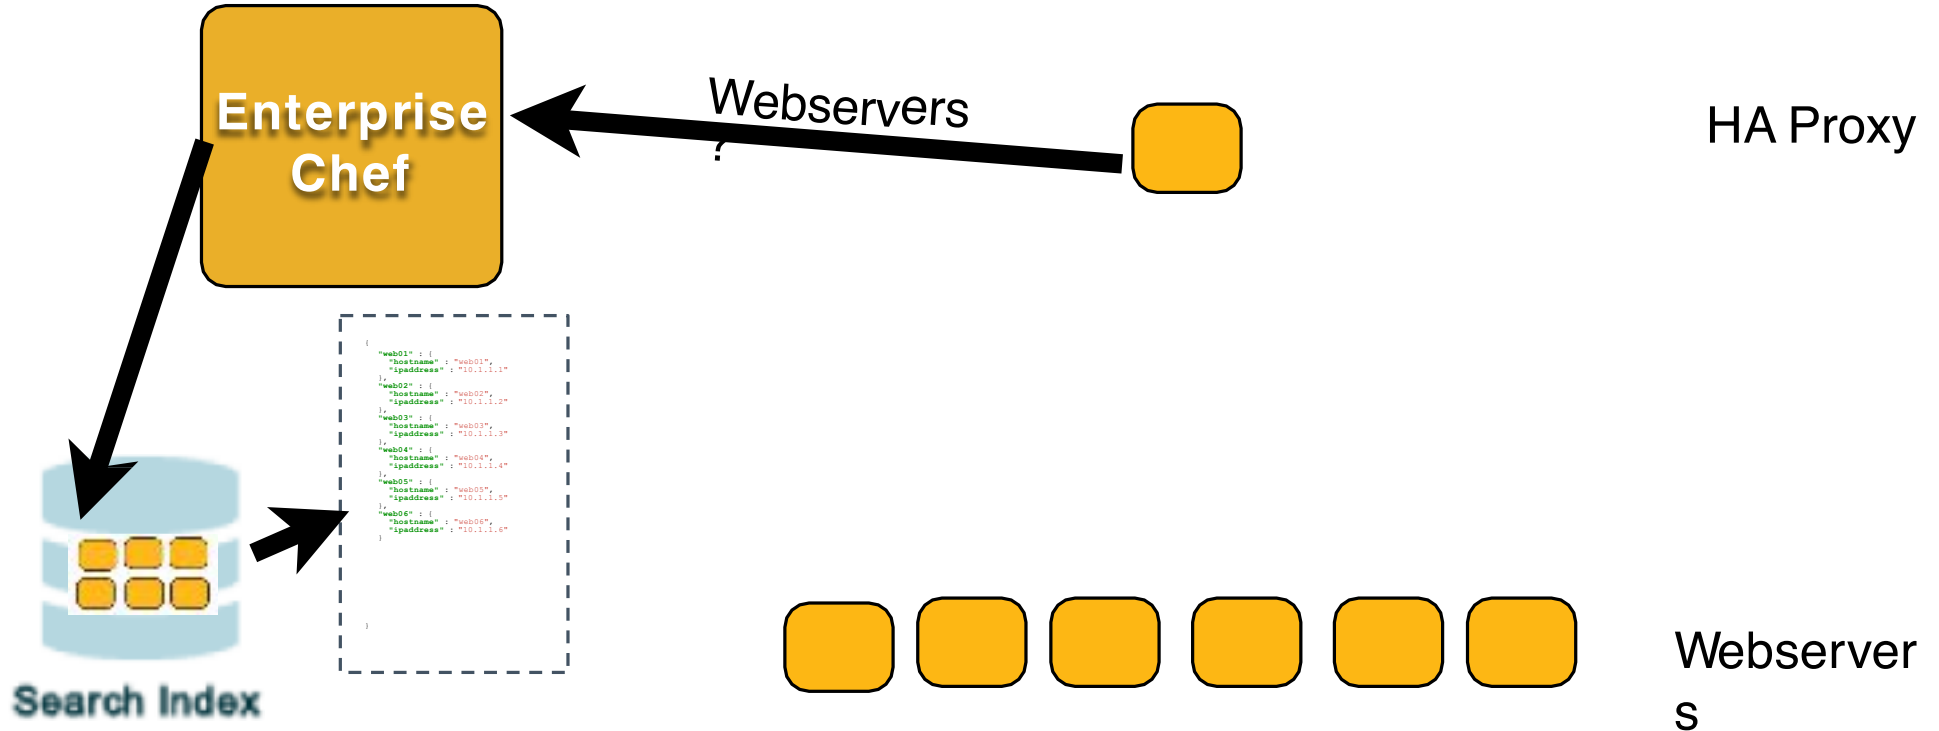
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



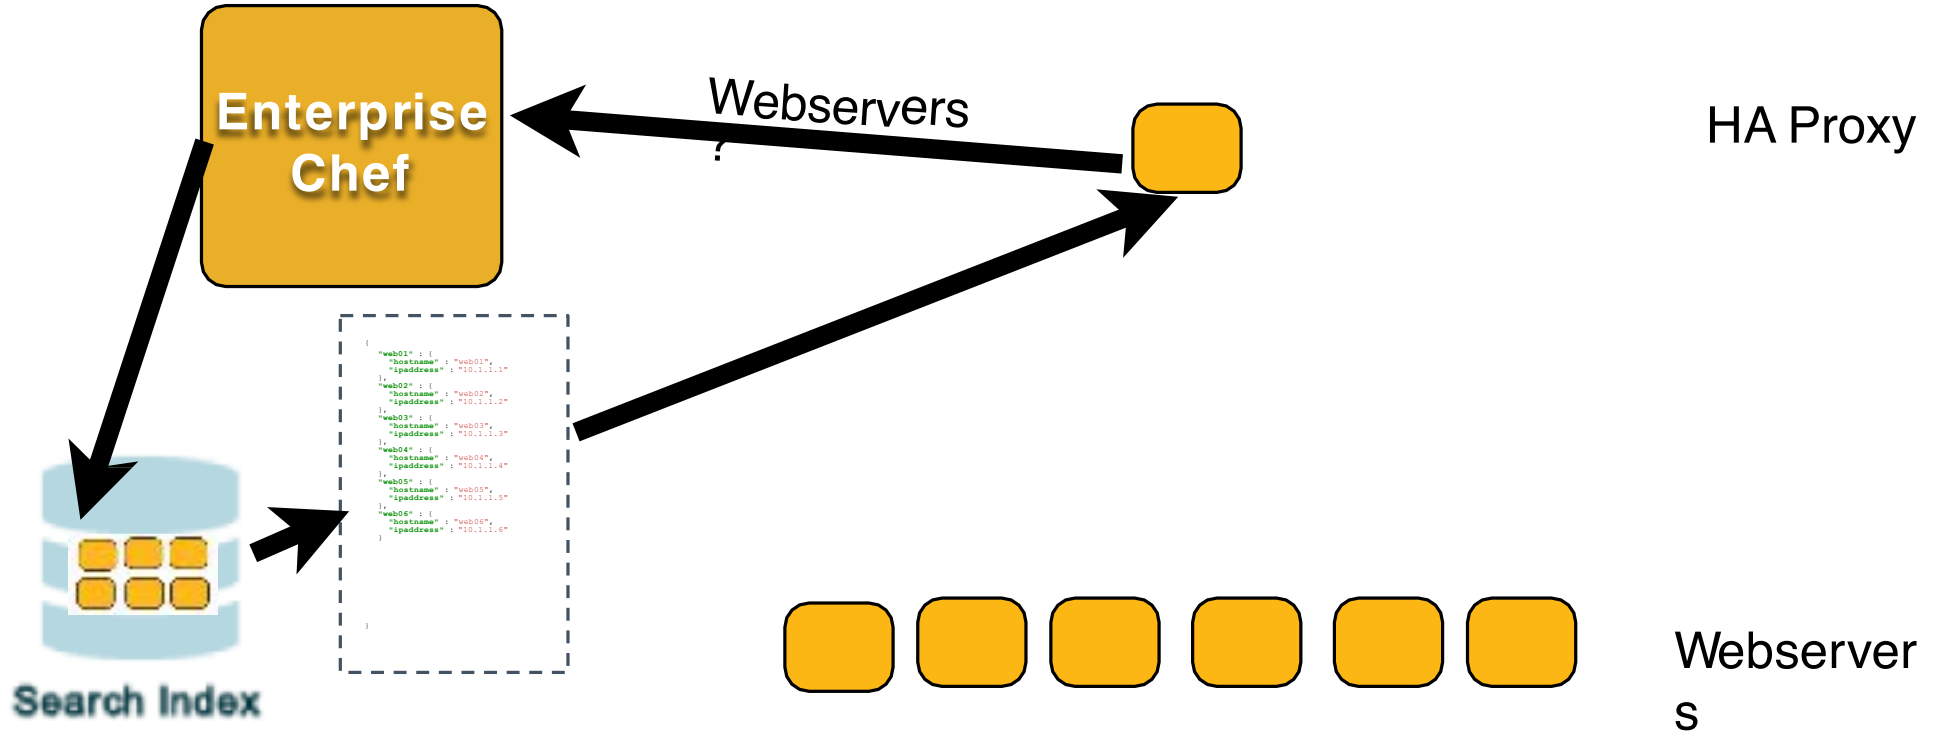
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



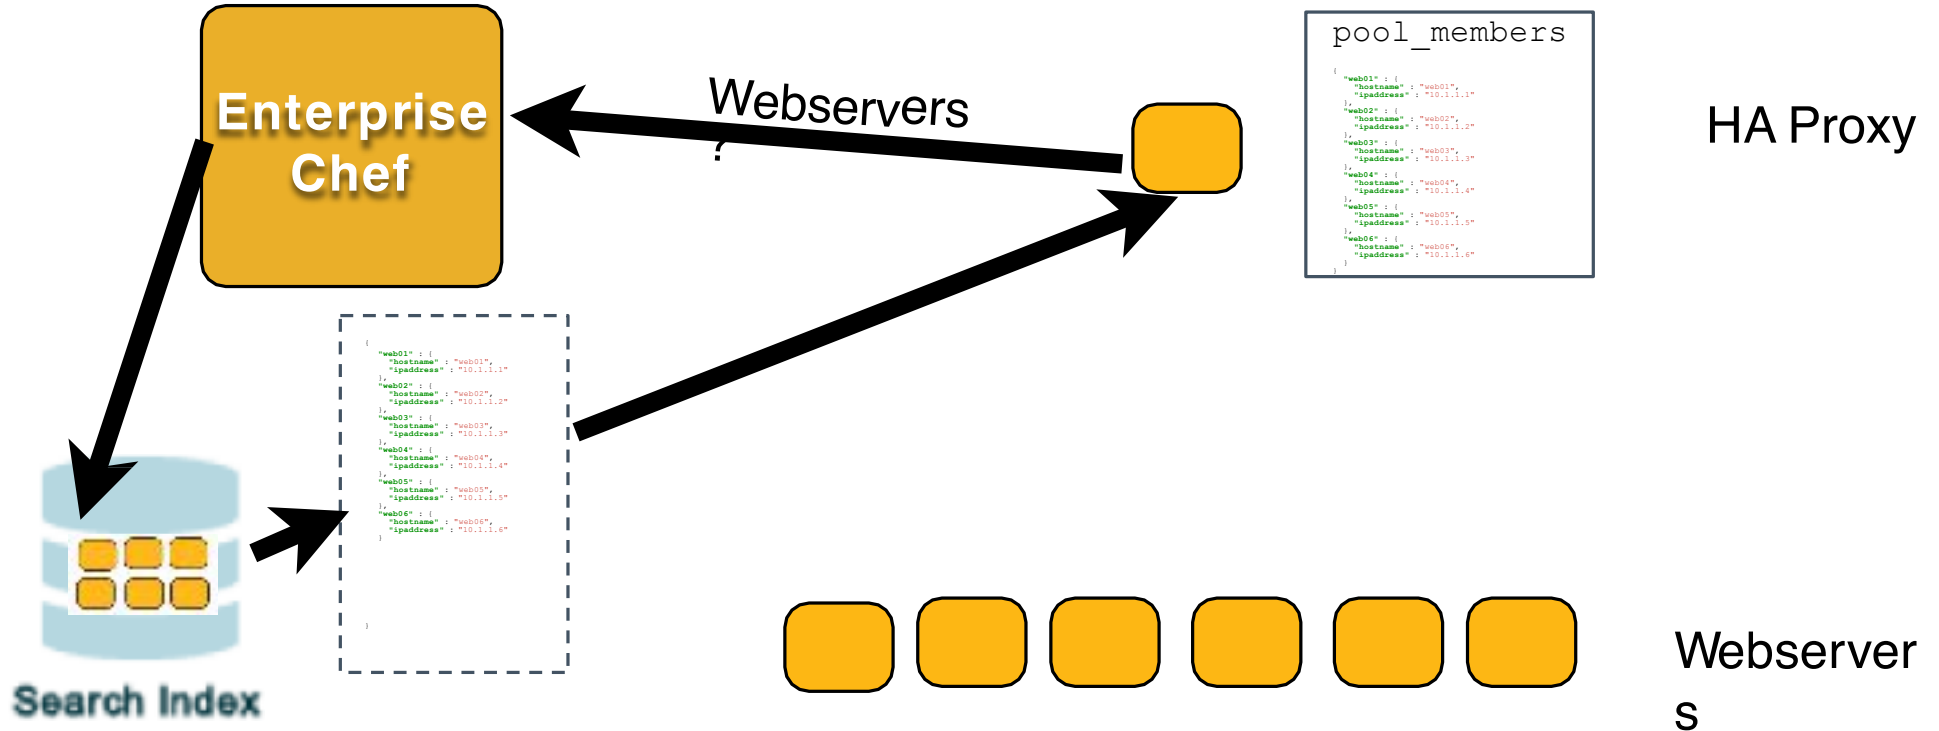
# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



# HAProxy Load Balancer

```
pool_members = search("node", "role:webserver")
```



# Search for Nodes

```
pool_members = search("node", "role:webserver")

template "/etc/haproxy/haproxy.cfg" do
  source "haproxy-app_lb.cfg.erb"
  owner "root"
  group "root"
  mode 0644
  variables :pool_members => pool_members.uniq

  notifies :restart, "service[haproxy]"
end
```

# Pass results into Templates

```
# Set up application listeners here.
listen application 0.0.0.0:80
  balance roundrobin
  <% @pool_members.each do |member| -%>

    server <%= member[:hostname] %> <%= member[:ipaddress] %>:>
weight 1 maxconn 1 check
  <% end -%>
<% if node["haproxy"]["enable_admin"] -%>
listen admin 0.0.0.0:22002
  mode http
  stats uri /
<% end -%>
```



# HAProxy Configuration

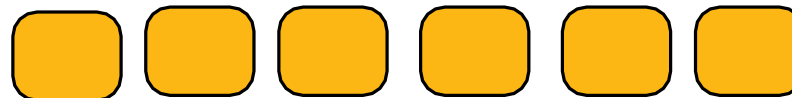
```
<% @pool_members.each do |member| -%>
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check
<% end -%>
```

pool\_members

```
{
  "web01" : {
    "hostname" : "web01",
    "ipaddress" : "10.1.1.1"
  },
  "web02" : {
    "hostname" : "web02",
    "ipaddress" : "10.1.1.2"
  },
  "web03" : {
    "hostname" : "web03",
    "ipaddress" : "10.1.1.3"
  },
  "web04" : {
    "hostname" : "web04",
    "ipaddress" : "10.1.1.4"
  },
  "web05" : {
    "hostname" : "web05",
    "ipaddress" : "10.1.1.5"
  },
  "web06" : {
    "hostname" : "web06",
    "ipaddress" : "10.1.1.6"
  }
}
```



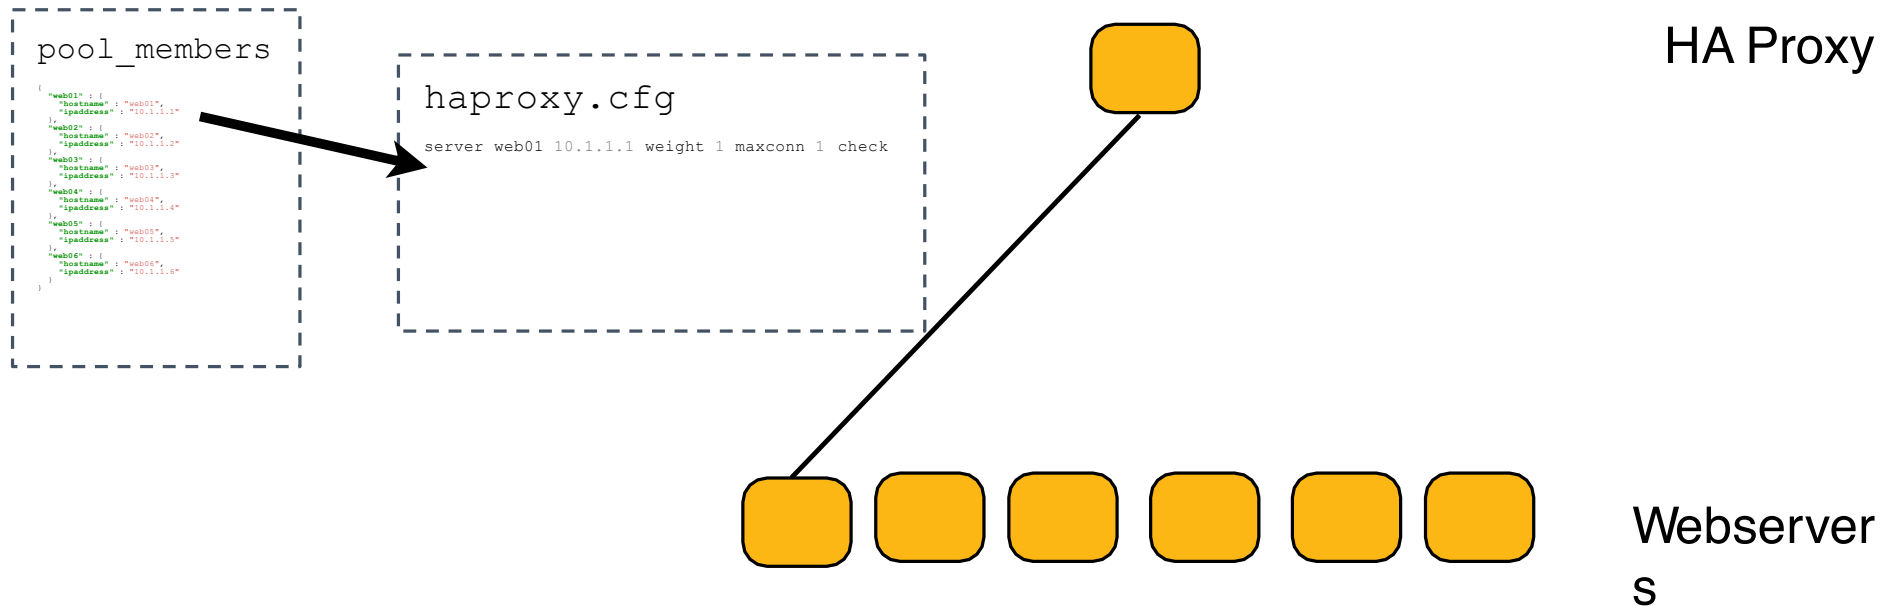
HA Proxy



Webserver  
s

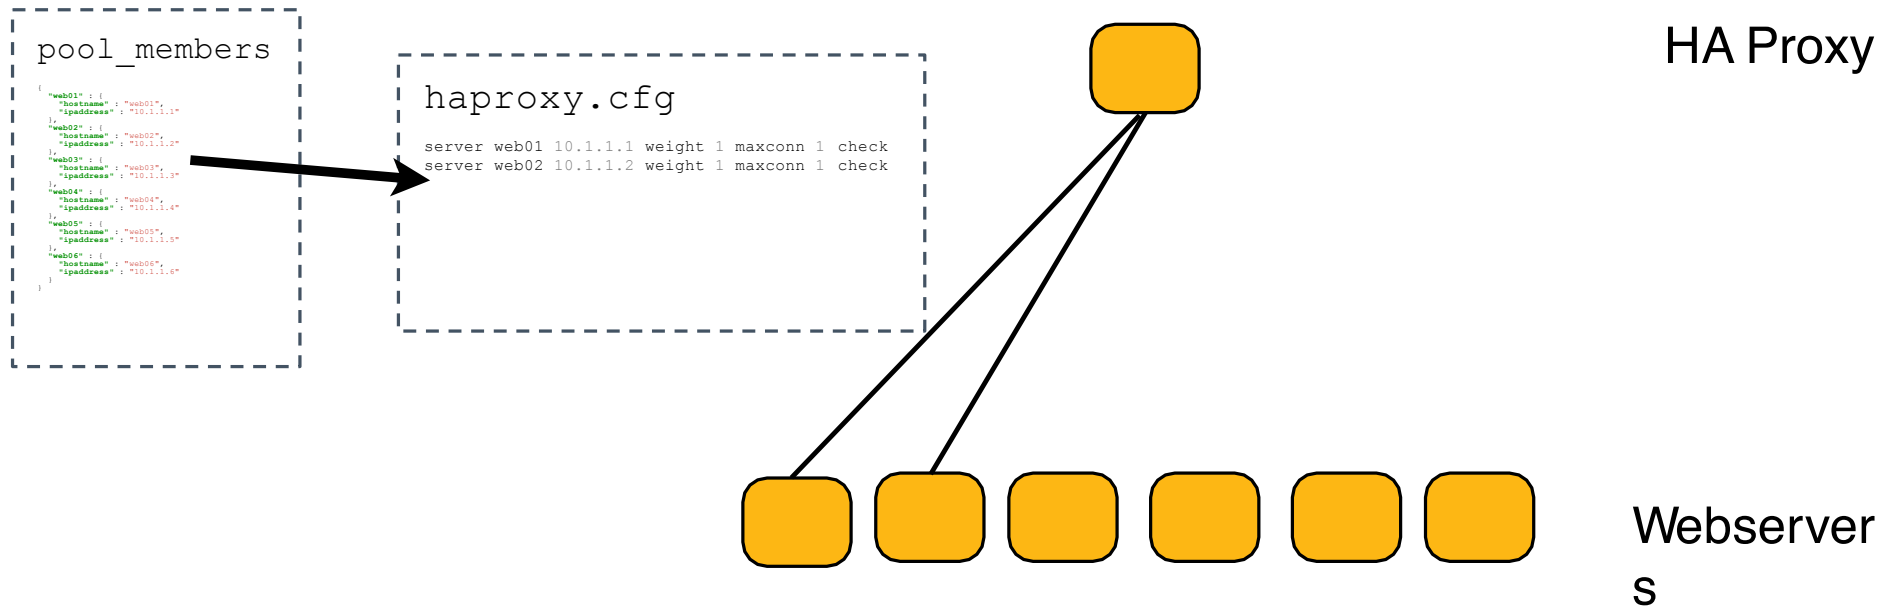
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check
<% end -%>
```



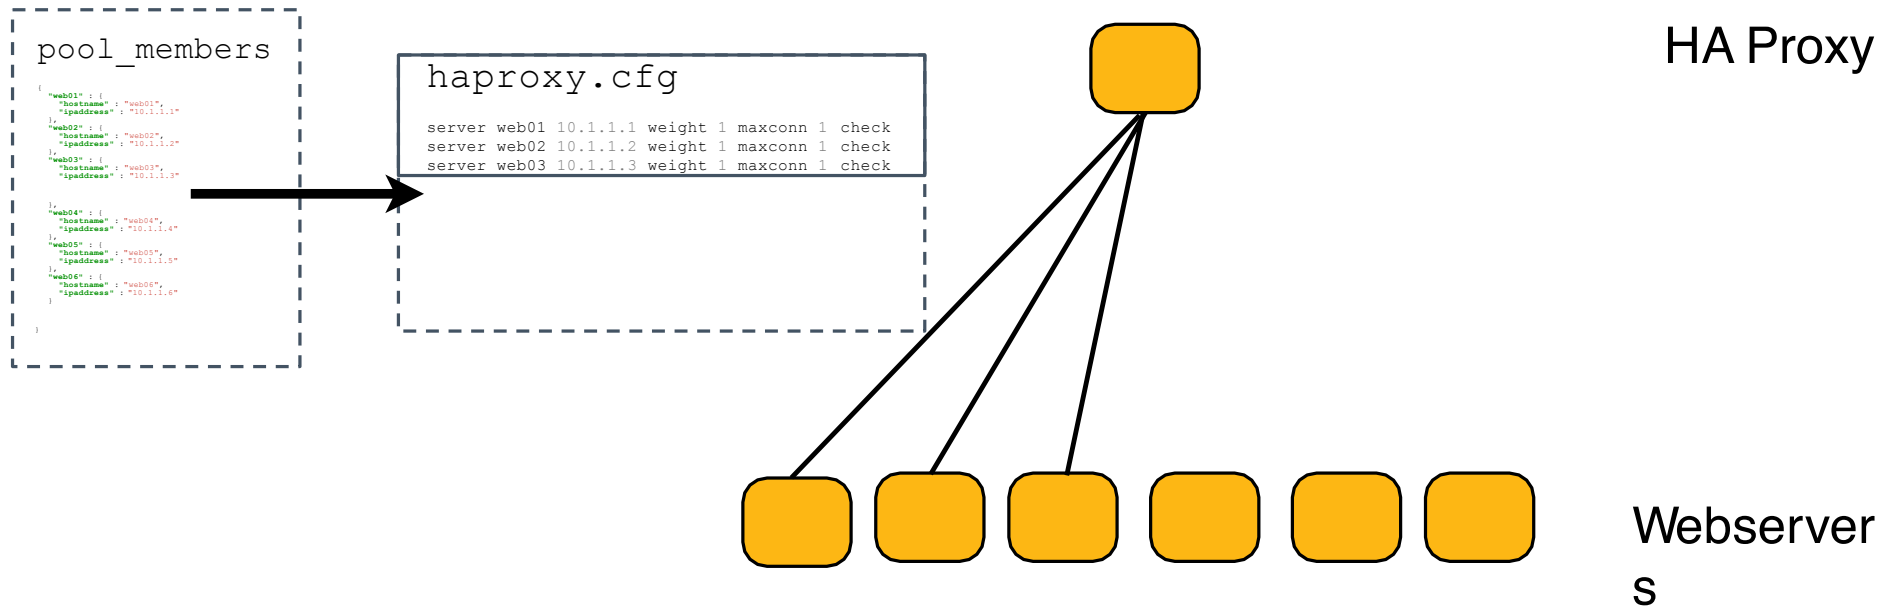
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check
<% end -%>
```



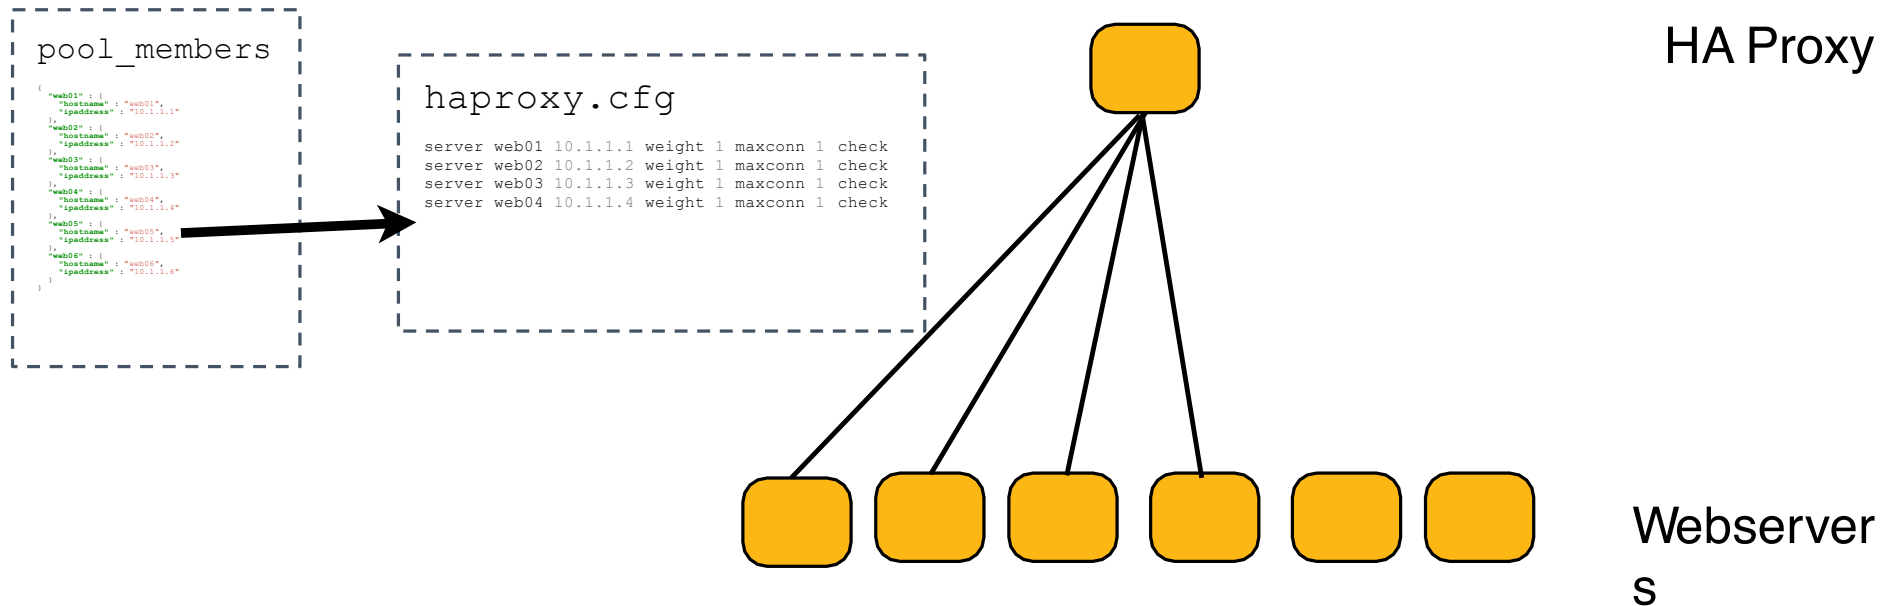
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member[:hostname] %> <%= member[:ipaddress] %>: weight 1 maxconn 1 check
<% end -%>
```



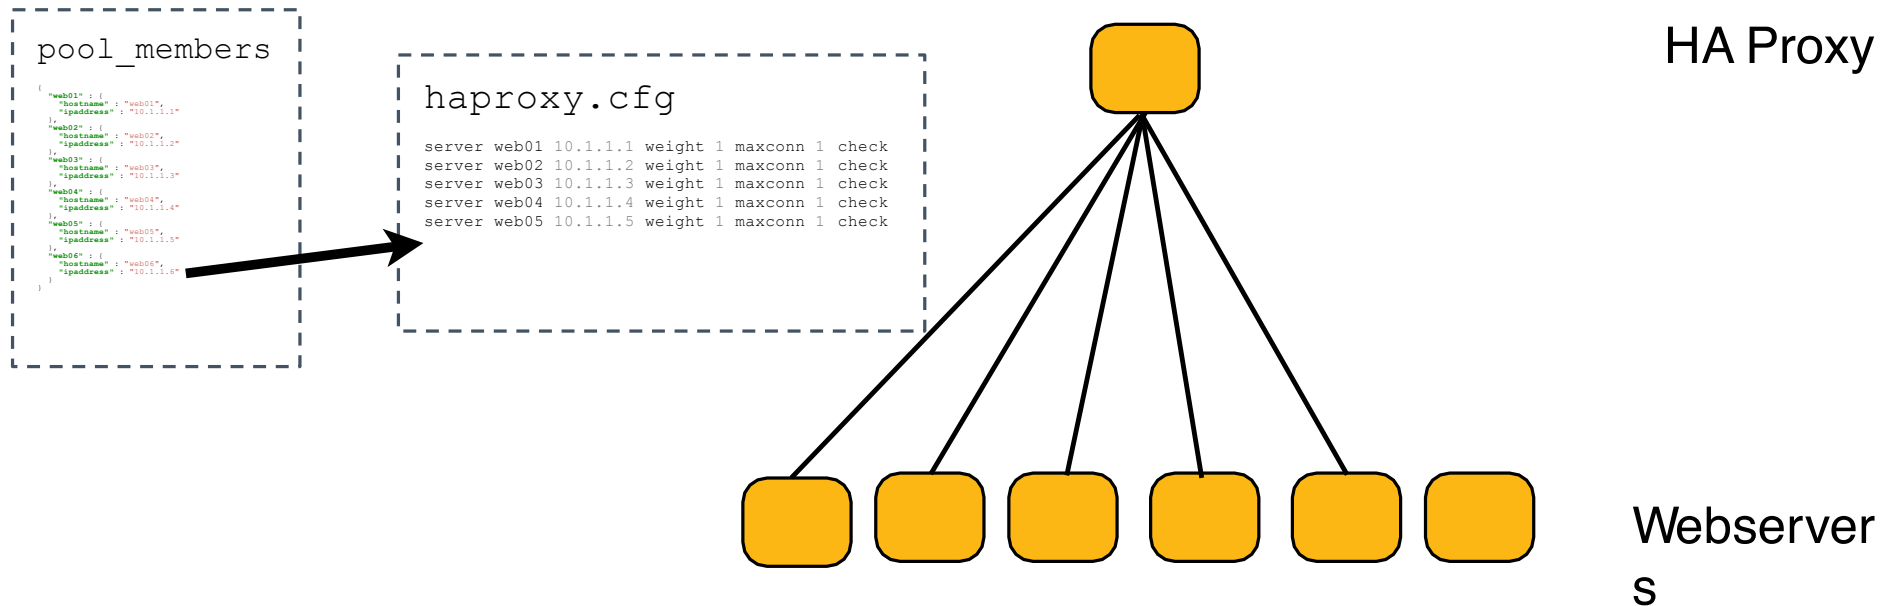
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>  
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check  
<% end -%>
```



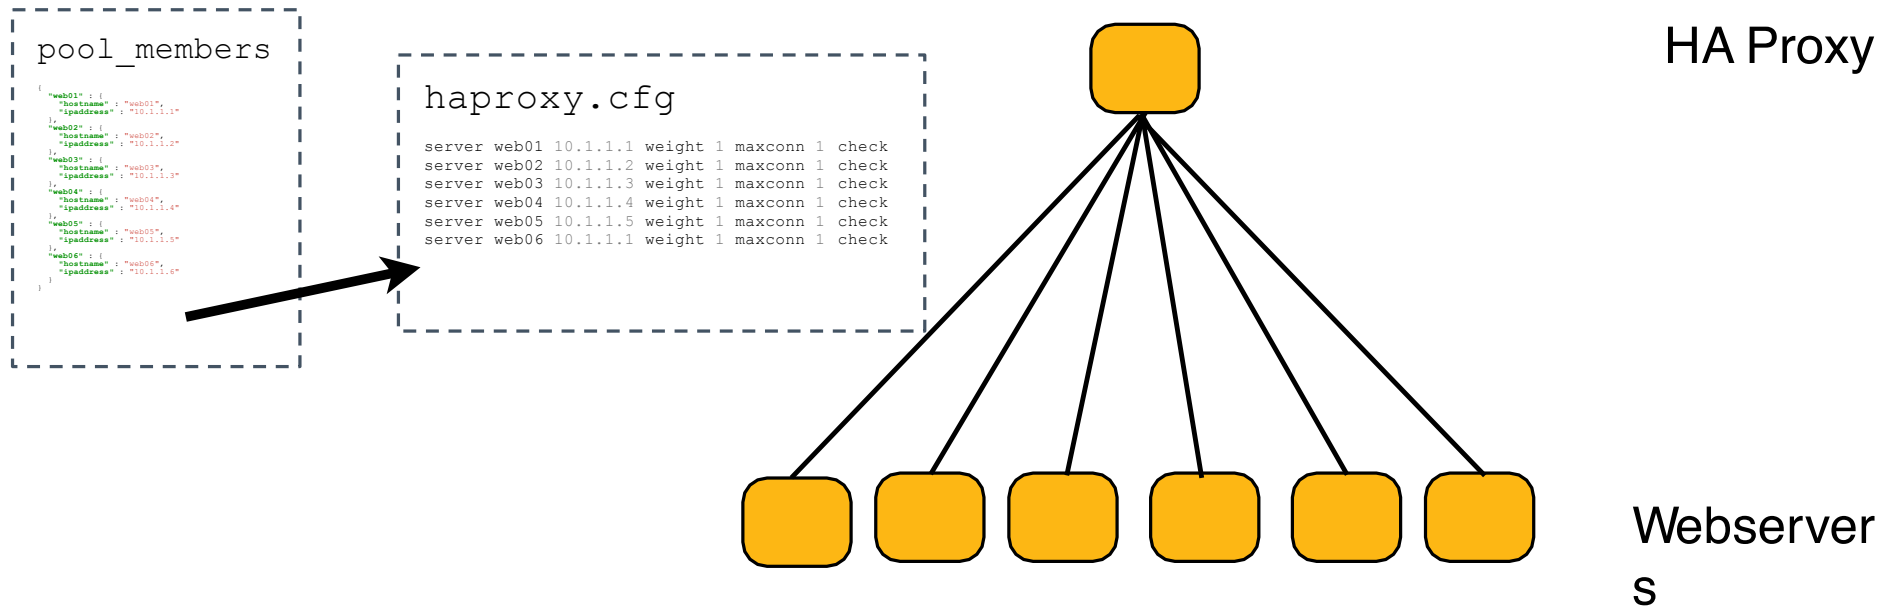
# HAProxy Configuration

```
<% @pool_members.each do |member| -%>  
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check  
<% end -%>
```

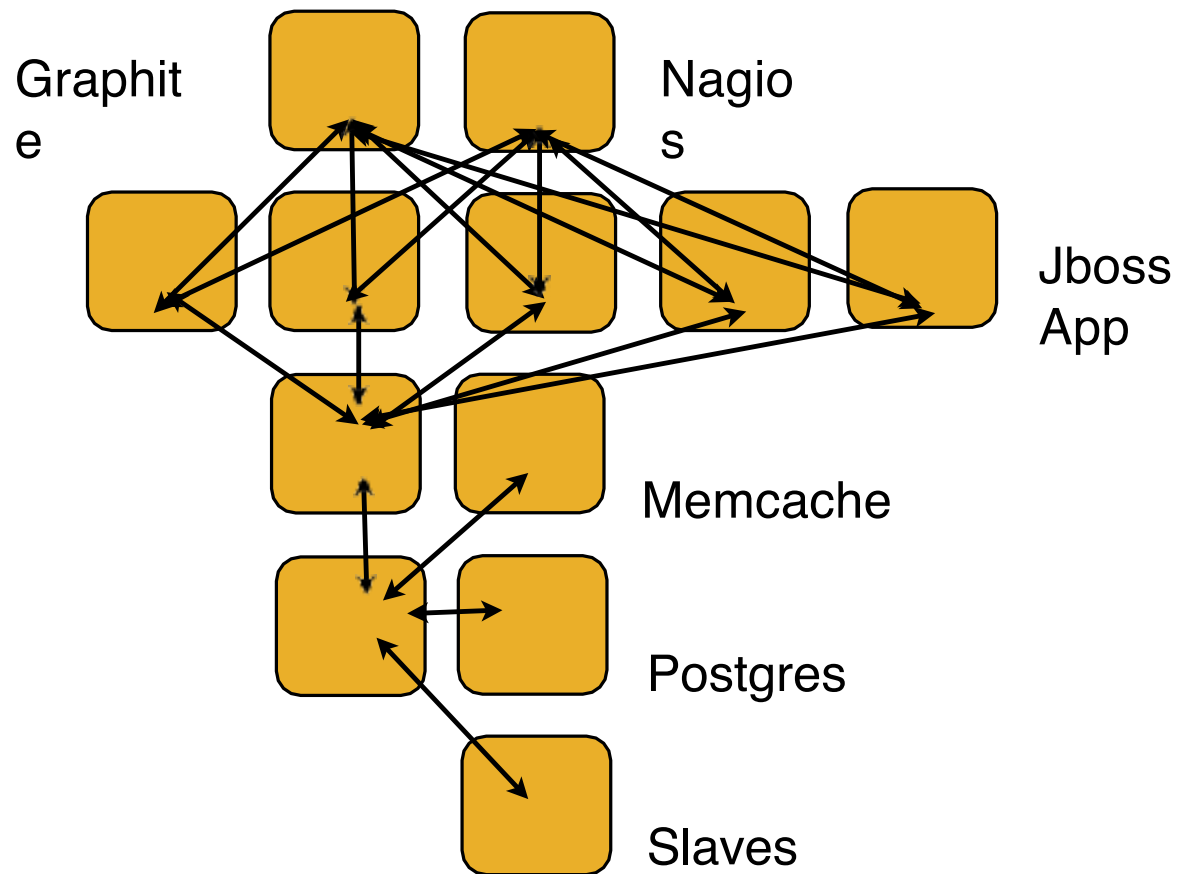


# HAProxy Configuration

```
<% @pool_members.each do |member| -%>
server <%= member[:hostname] %> <%= member[:ipaddress] %>:> weight 1 maxconn 1 check
<% end -%>
```

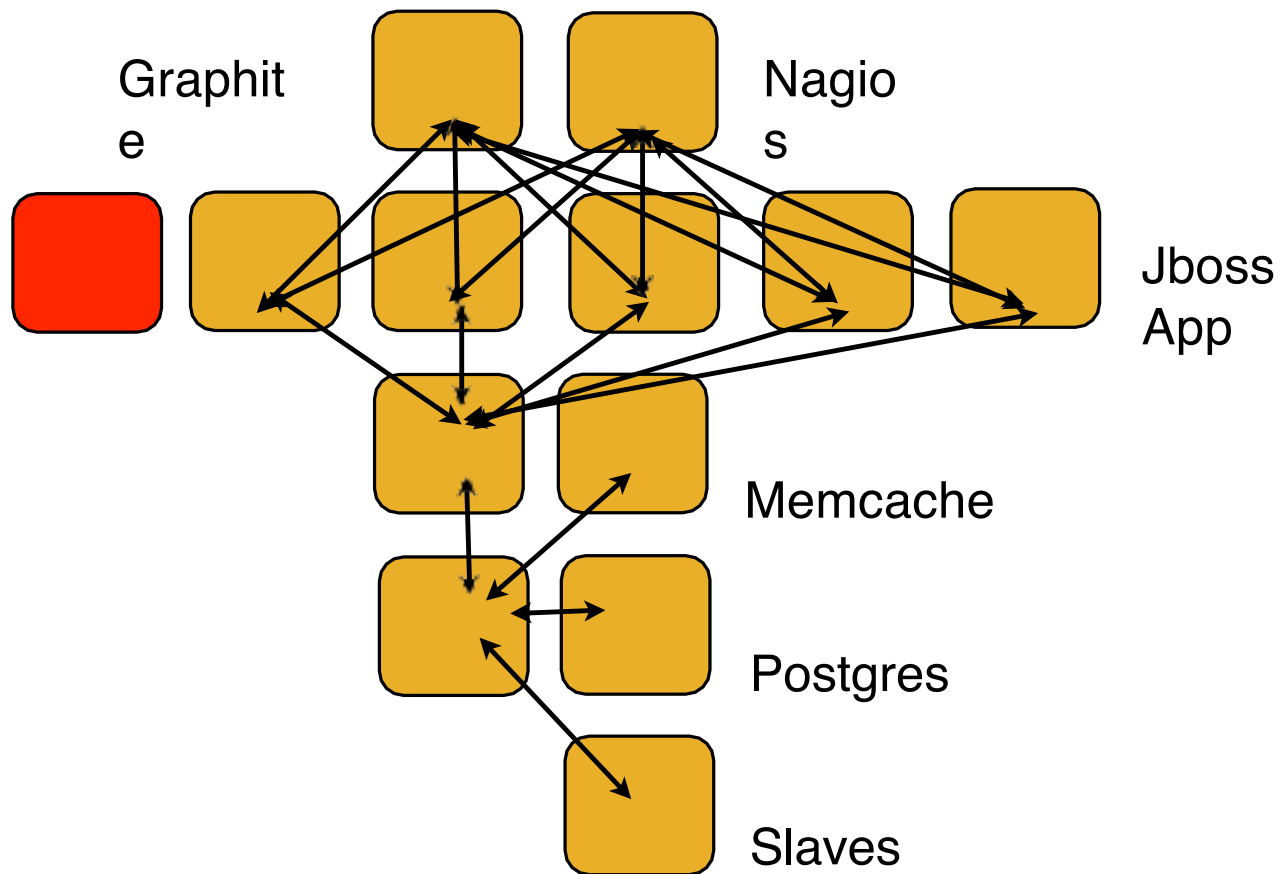


# So when this....

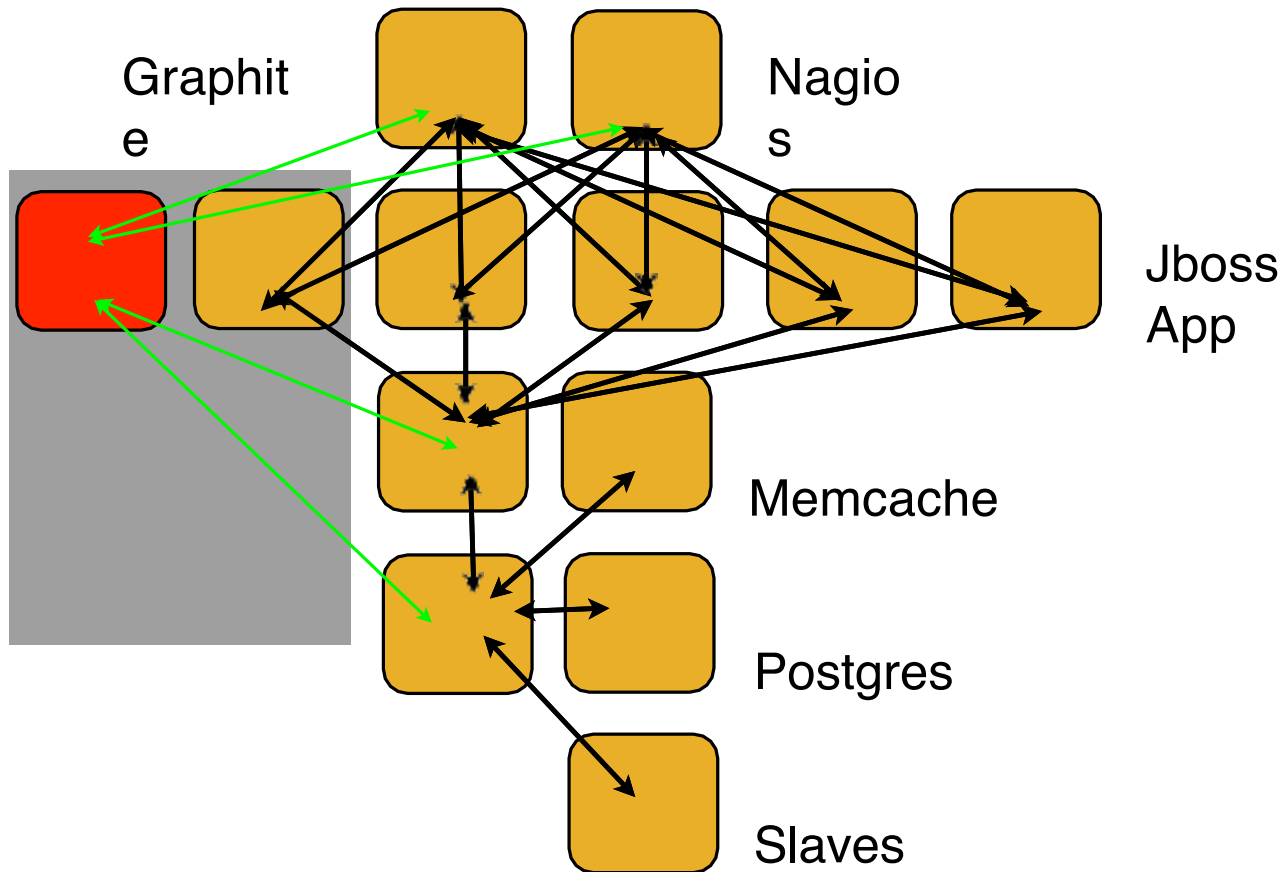




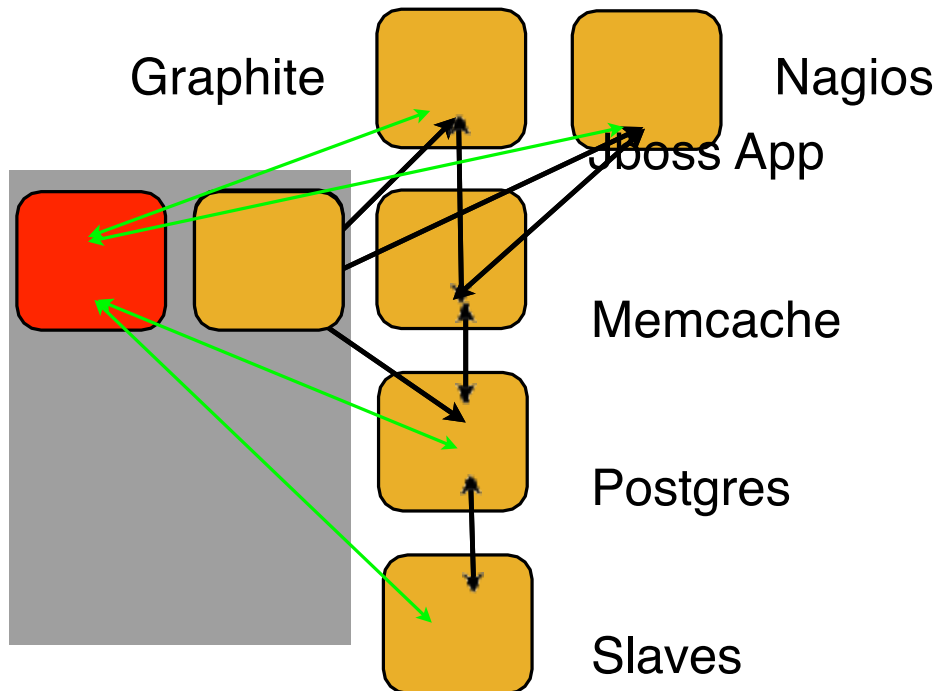
# ...becomes this



# ...this can happen automatically



# Count the Resources



- Load balancer config
- Nagios host ping
- Nagios host ssh
- Nagios host HTTP
- Nagios host app health
- Graphite CPU
- Graphite Memory
- Graphite Disk
- Graphite SNMP
- Memcache firewall
- Postgres firewall
- Postgres authZ config

# Manage Complexity

- Determine the desired state of your infrastructure
- Identify the Resources required to meet that state
- Gather the Resources into Recipes
- Compose a Run List from Recipes
- Apply a Run List to each Node in your environment
- Your infrastructure adheres to the policy modeled in Chef

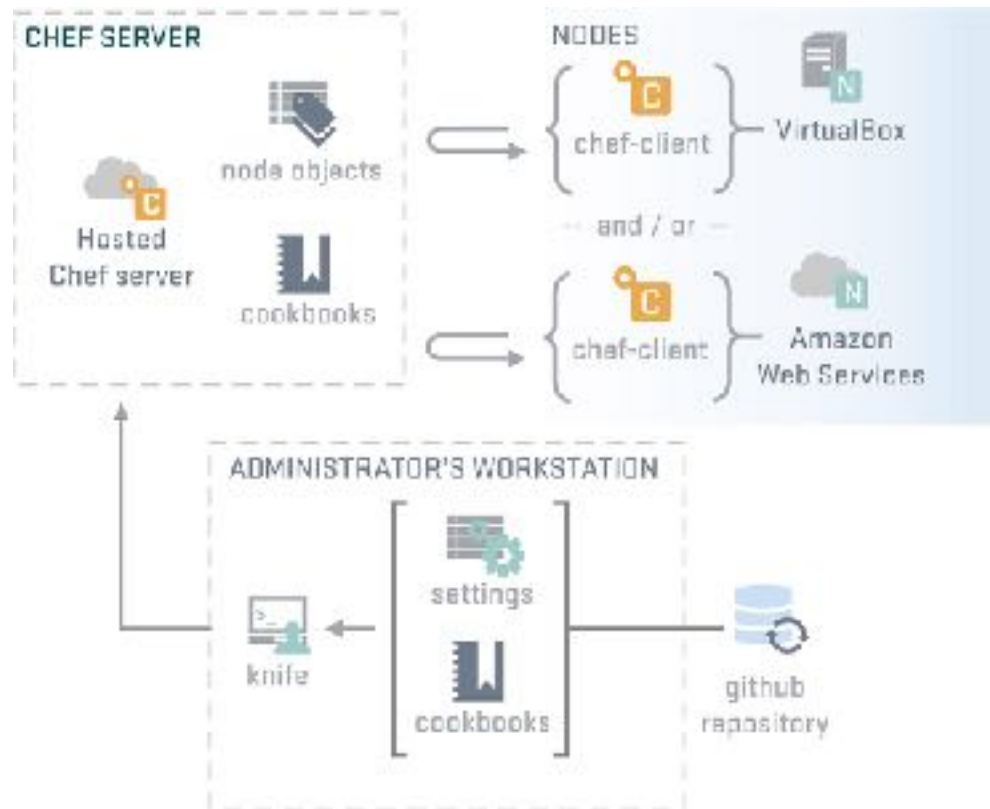
# Configuration Drift

- Configuration Drift happens when:
  - Your infrastructure requirements change
  - The configuration of a server falls out of policy
- Chef makes it easy to manage
  - Model the new requirements in your Chef configuration files
  - Run the chef-client to enforce your policies

# Recap

- In this section, we have
  - Described how Chef thinks about Infrastructure Automation
  - Defined the following terms:
    - Resource
    - Recipe
    - Node
    - Run List
    - Search

# Chef Infrastructure



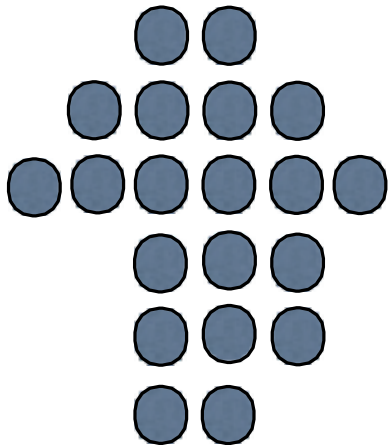
# Sign-up for Hosted Chef

- <http://getchef.com>
- Click “Get Chef”
- Select “Hosted Chef”
- Complete the registration form
- Create an Organization

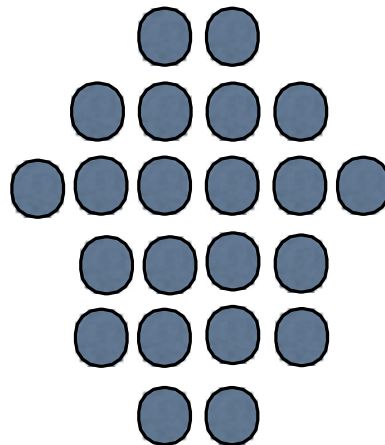


# Organizations

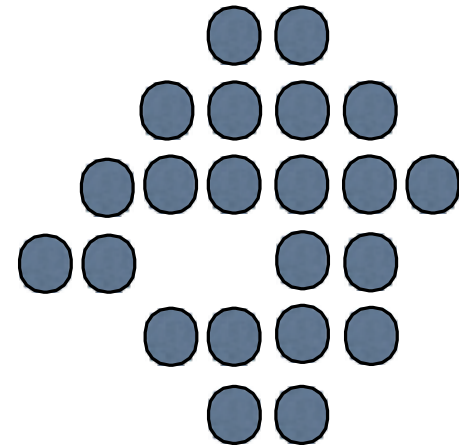
My  
Infrastructure



Your  
Infrastructure



Their  
Infrastructure



# Organizations

- Provide multi-tenancy in Enterprise Chef
- Nothing is shared between Organizations - they're completely independent
- May represent different
  - Companies
  - Business Units
  - Departments

# Configure Workstation

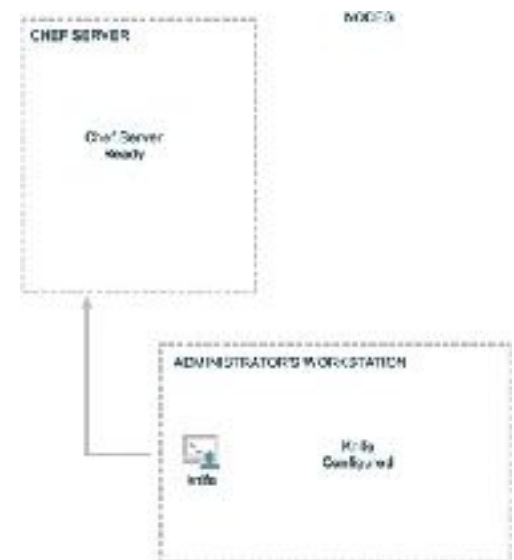
- Download and extract Chef starter kit
- Install chef-client
  - <http://getchef.com/chef/install>

# Verify Knife

```
$ knife --version  
Chef: 11.12.4
```

- Your version may be different, that's ok!

```
$ knife client list  
ORGNAME-validator
```



# knife.rb

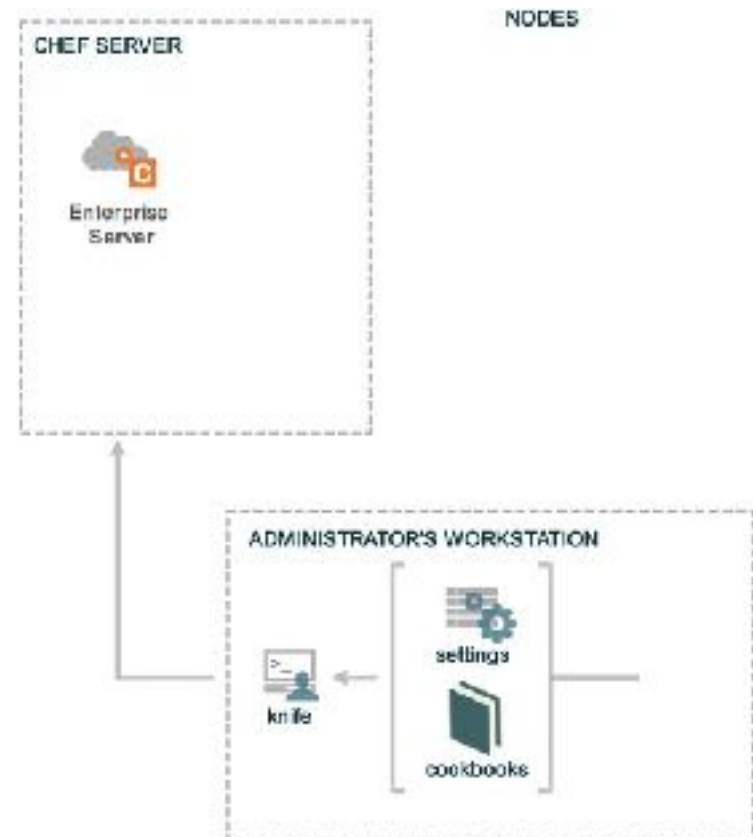


**OPEN IN EDITOR:** chef-repo/.chef/knife.rb

```
current_dir = File.dirname(__FILE__)
log_level      :info
log_location   STDOUT
node_name      "USERNAME"
client_key     "#{current_dir}/USERNAME.pem"
validation_client_name "ORGNAME-validator" "#{current_dir}/
validation_key  ORGNAME-validator.pem"
chef_server_url "https://api.opscode.com/organizations/ORGNAME"
cache_type     'BasicFile'
cache_options( :path => "#{ENV['HOME']}/.chef/checksums" )
cookbook_path  ["#{current_dir}/../cookbooks"]
```

# knife client list

1. Reads the `chef_server_url` from `knife.rb`
2. Invokes HTTP GET to `#{chef_server_url}/clients`
3. Displays the result



# Additional Resources

- Chef Fundamentals Webinar Series
- <https://www.youtube.com/watch?v=S5lHUpzoCYo&list=PL11cZfNdWNyPnZA9D1MbVqldGuOWqbumZ>
- Discussion group for webinar participants
- <https://groups.google.com/d/forum/learnchef-fundamentals-webinar>

# Additional Resources

- Learn Chef
- <http://learnchef.com>
- Documentation
- <http://docs.opscode.com>
- Slides are from [chef.com](http://chef.com)